

Deep Dive on Amazon Redshift



Agenda

Deep inside Redshift Architecture

10%

Performance tuning

5%

Integration with AWS data services

5%

Redshift Spectrum

10%

Redshift Echo System

15%

Redshift at Manchester Airport Group

45%

Summary + Q&A

10%

Redshift Architecture



Fast



**Cost
Efficient**



Simple



Elastic



Secure



Compatible



Amazon Redshift

Fast, simple, cost-effective data warehousing.

Fast, simple, cost-effective data warehousing.

*Managed Massively Parallel Petabyte Scale Data
Warehouse*

Streaming Backup/Restore to S3

Load data from S3, DynamoDB and EMR

Extensive Security Features

Scale from 160 GB -> 2 PB Online

Amazon Redshift Cluster Architecture

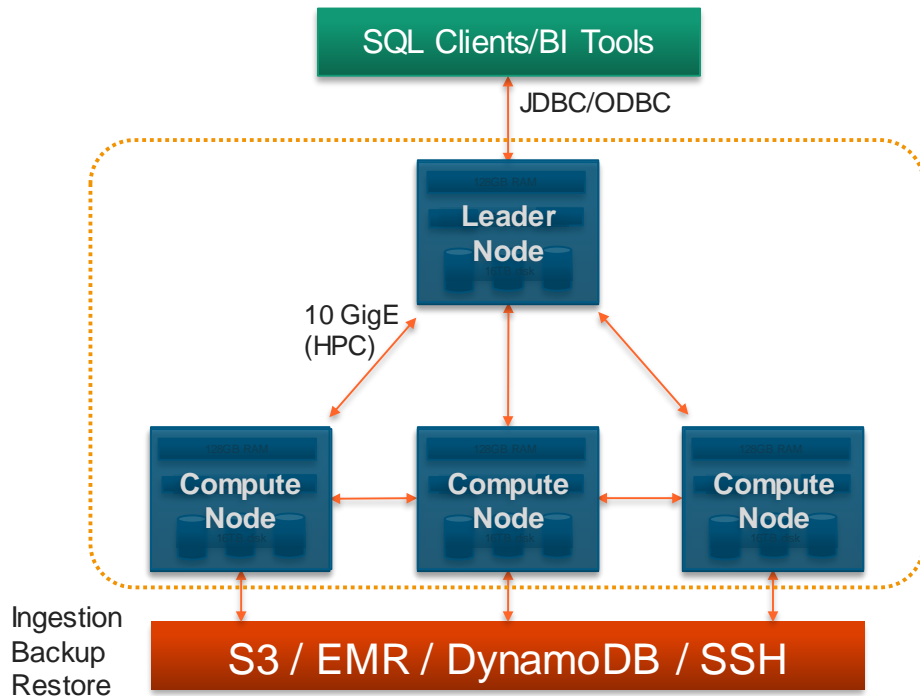
Massively parallel, shared nothing

Leader node

- SQL endpoint
- Stores metadata
- Coordinates parallel SQL processing

Compute nodes

- Local, columnar storage
- Executes queries in parallel
- Load, backup, restore
- 2, 16 or 32 slices



Design for Queryability


- *Equally* on each slice
- *Minimum* amount of work
- Use *just enough* cluster resources

Do an *Equal* Amount of Work
on Each Slice

Choose Best Table Distribution Style


Key

*Same key to
same location*




Even

*Round robin
distribution*



All

*All data on
every node*




**Do the *Minimum* Amount of
Work on Each Slice**

Reduced I/O = Enhanced Performance

Columnar storage
+
Large data block sizes
+
Data compression
+
Zone maps
+
Direct-attached storage

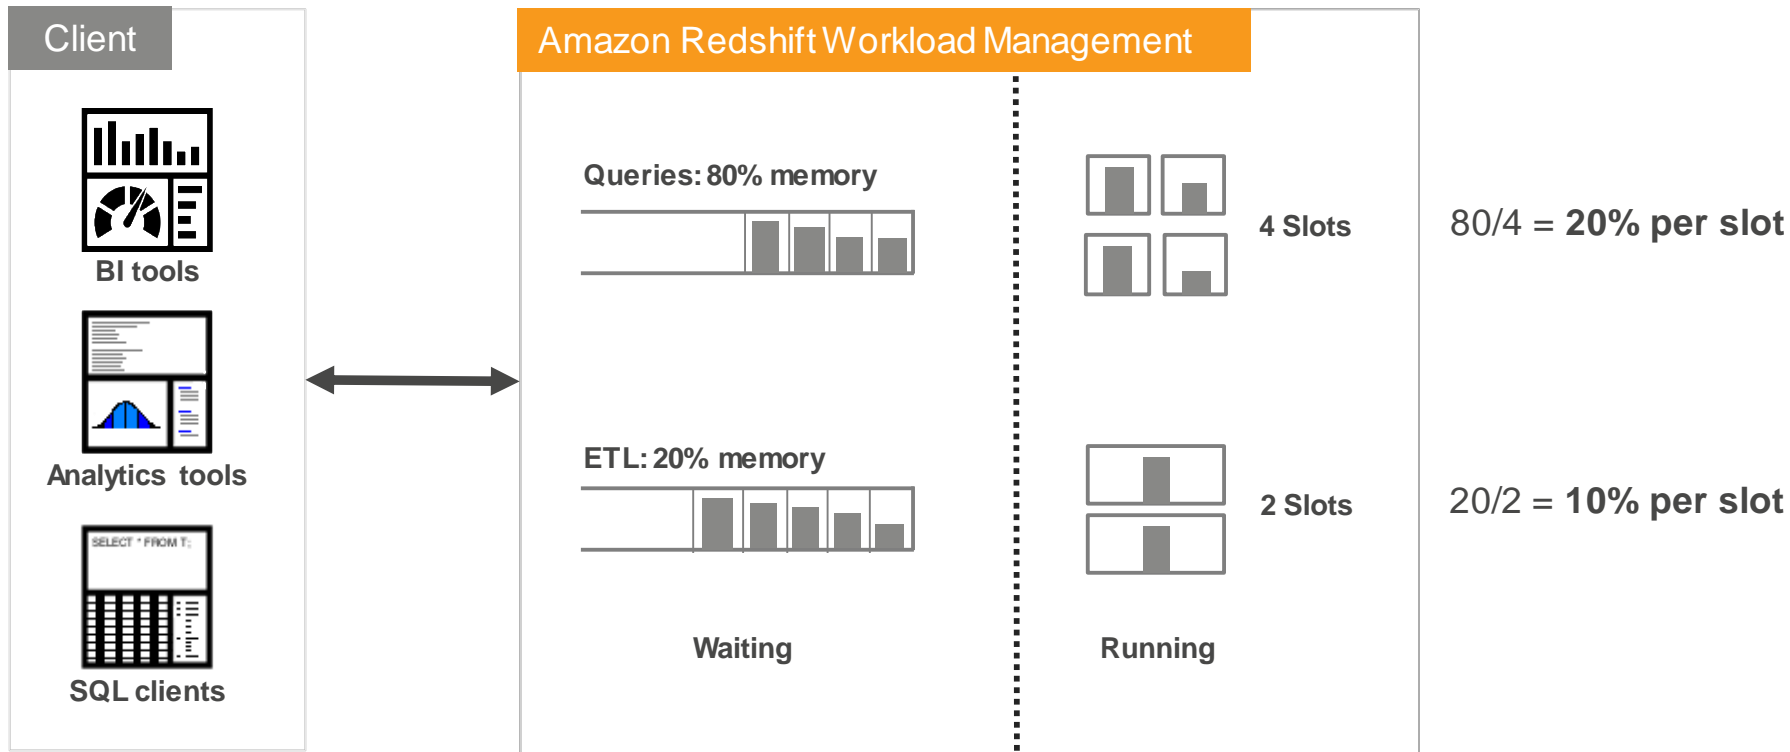
```
analyze compression listing;
```

Table	Column	Encoding
listing	listid	delta
listing	sellerid	delta32k
listing	eventid	delta32k
listing	dateid	bytedict
listing	numtickets	bytedict
listing	priceperticket	delta32k
listing	totalprice	mostly32
listing	listtime	raw



**Use Cluster Resources
Efficiently to Complete as
Quickly as Possible**

Workload Management



Architecture

Tuning

Integration

Spectrum

Echo
System

MAG

Summary

Redshift Performance Tuning

Redshift Playbook

Part 1: Preamble, Prerequisites, and
Prioritization

Part 2: **Distribution** Styles and
Distribution Keys

Part 3: Compound and Interleaved
Sort Keys

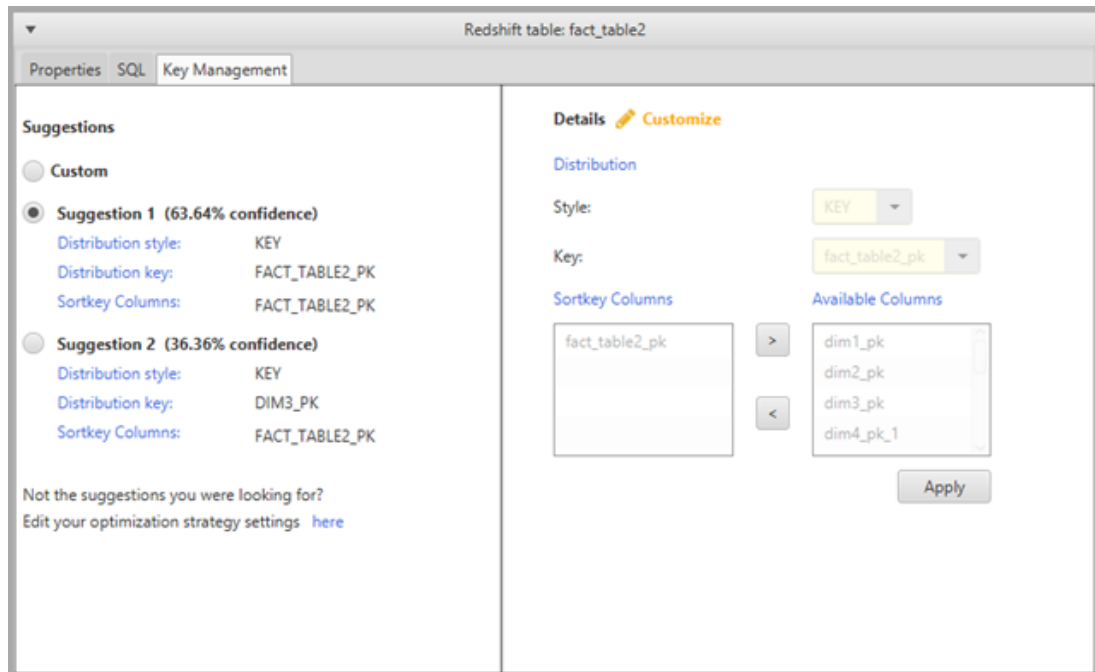
Part 4: Compression **Encodings**

Part 5: Table Data **Durability**



amzn.to/2quChdM

Optimizing Amazon Redshift by Using the AWS Schema Conversion Tool



amzn.to/2sTYow1

Architecture

Tuning

Integration

Spectrum

Echo
System

MAG

Summary

Ingestion, ETL & BI

Getting data to Redshift using AWS DMS



Simple to use



Minimal Downtime



Supports most widely used Databases



Low Cost



Fast & Easy to Set-up



Reliable

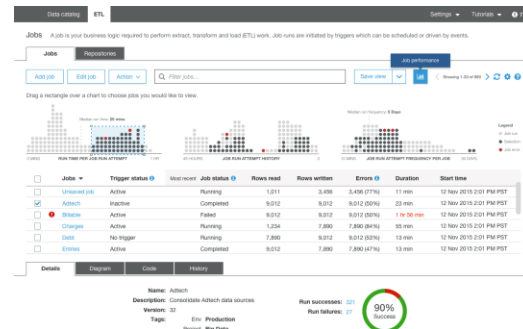
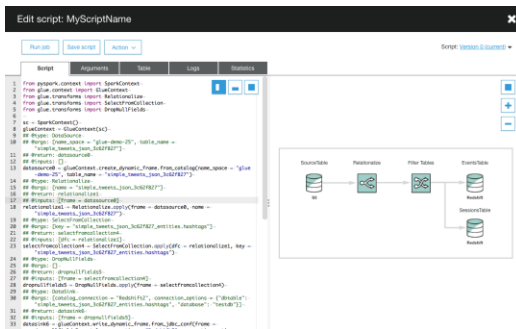
Loading data from S3

- Splitting Your Data into Multiple Files
- Uploading Files to Amazon S3
- Using the COPY Command to Load from Amazon S3

Fully-managed data catalog and ETL service



S3, RDS, Redshift & any
JDBC-compliant data store




QuickSight for BI on Redshift

QuickSight seamlessly connects to Redshift giving you native access to all of your clusters and tables.



Amazon Redshift



Achieve high concurrency by offloading end user queries to SPICE.

Calculations can be done in SPICE reducing the load on the underlying database.



Architecture

Tuning

Integration

Spectrum

Echo
System

MAG

Summary


Amazon Redshift Spectrum

Amazon Redshift Spectrum

Run SQL queries directly against data in S3 using thousands of nodes




Fast @ exabyte scale



Elastic & highly available




On-demand, pay-per-query



High concurrency: Multiple clusters access same data

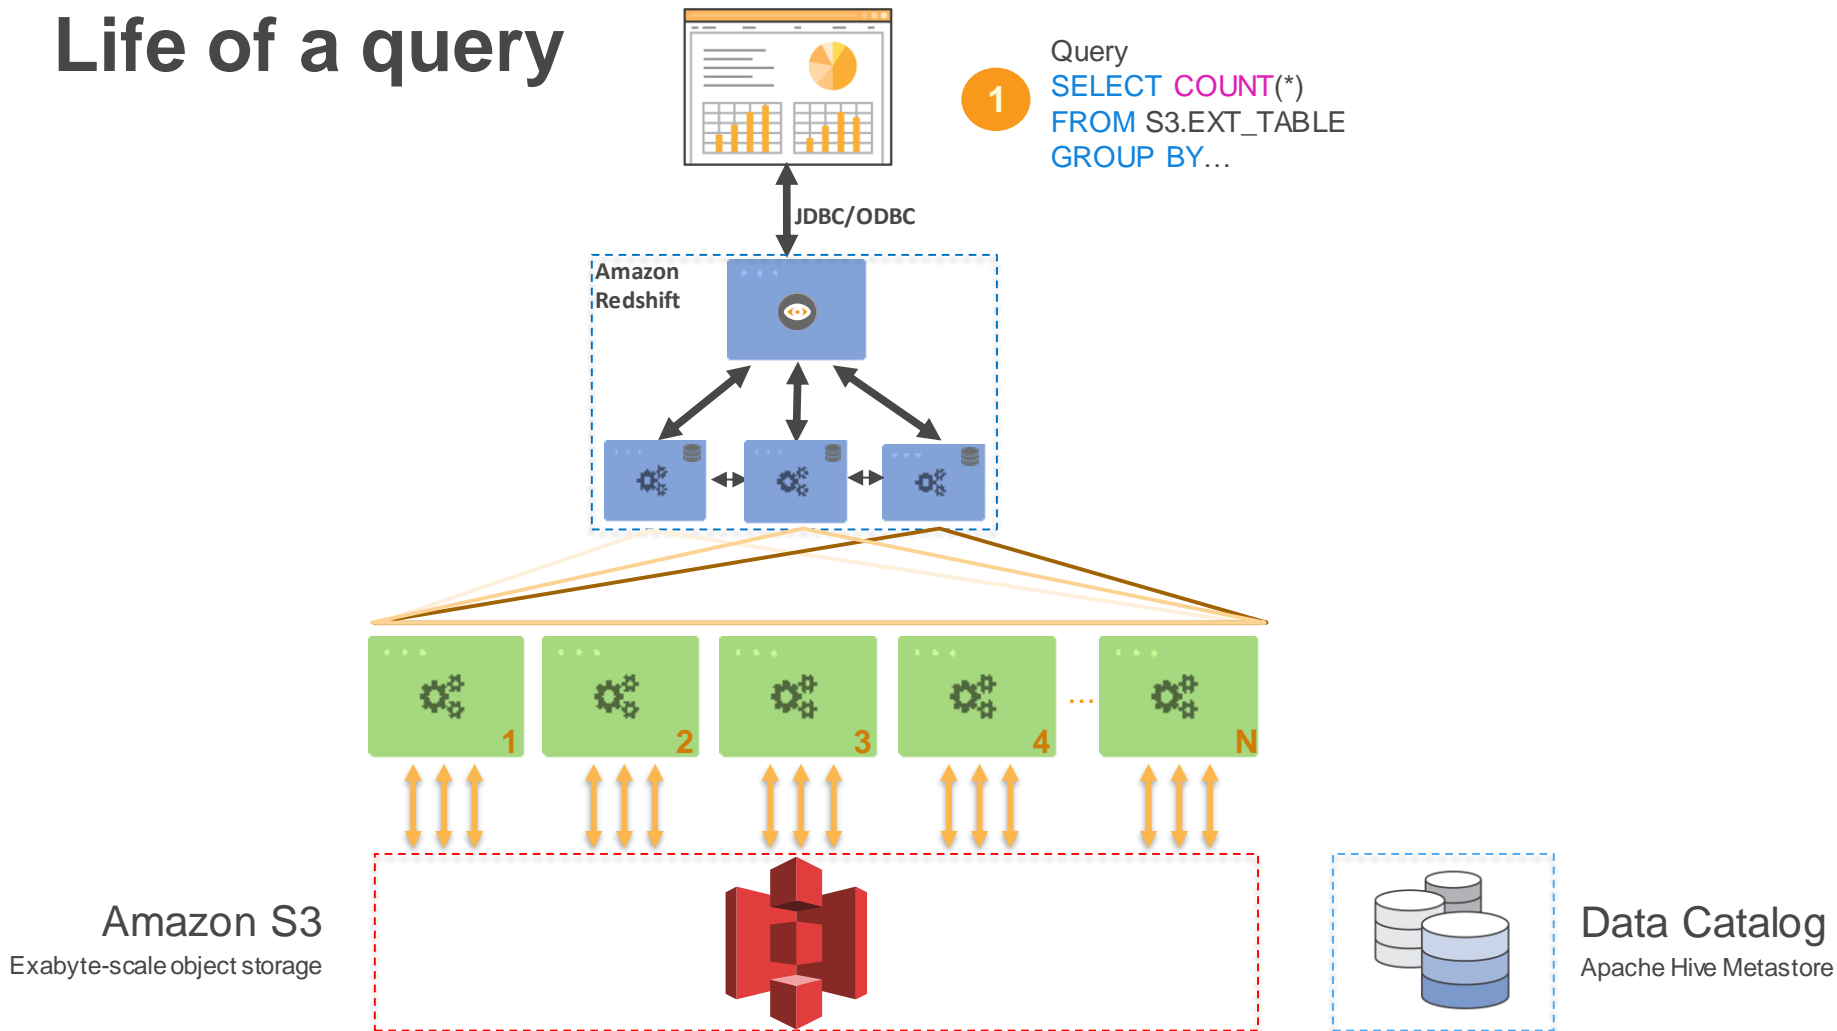


No ETL: Query data in-place using open file formats

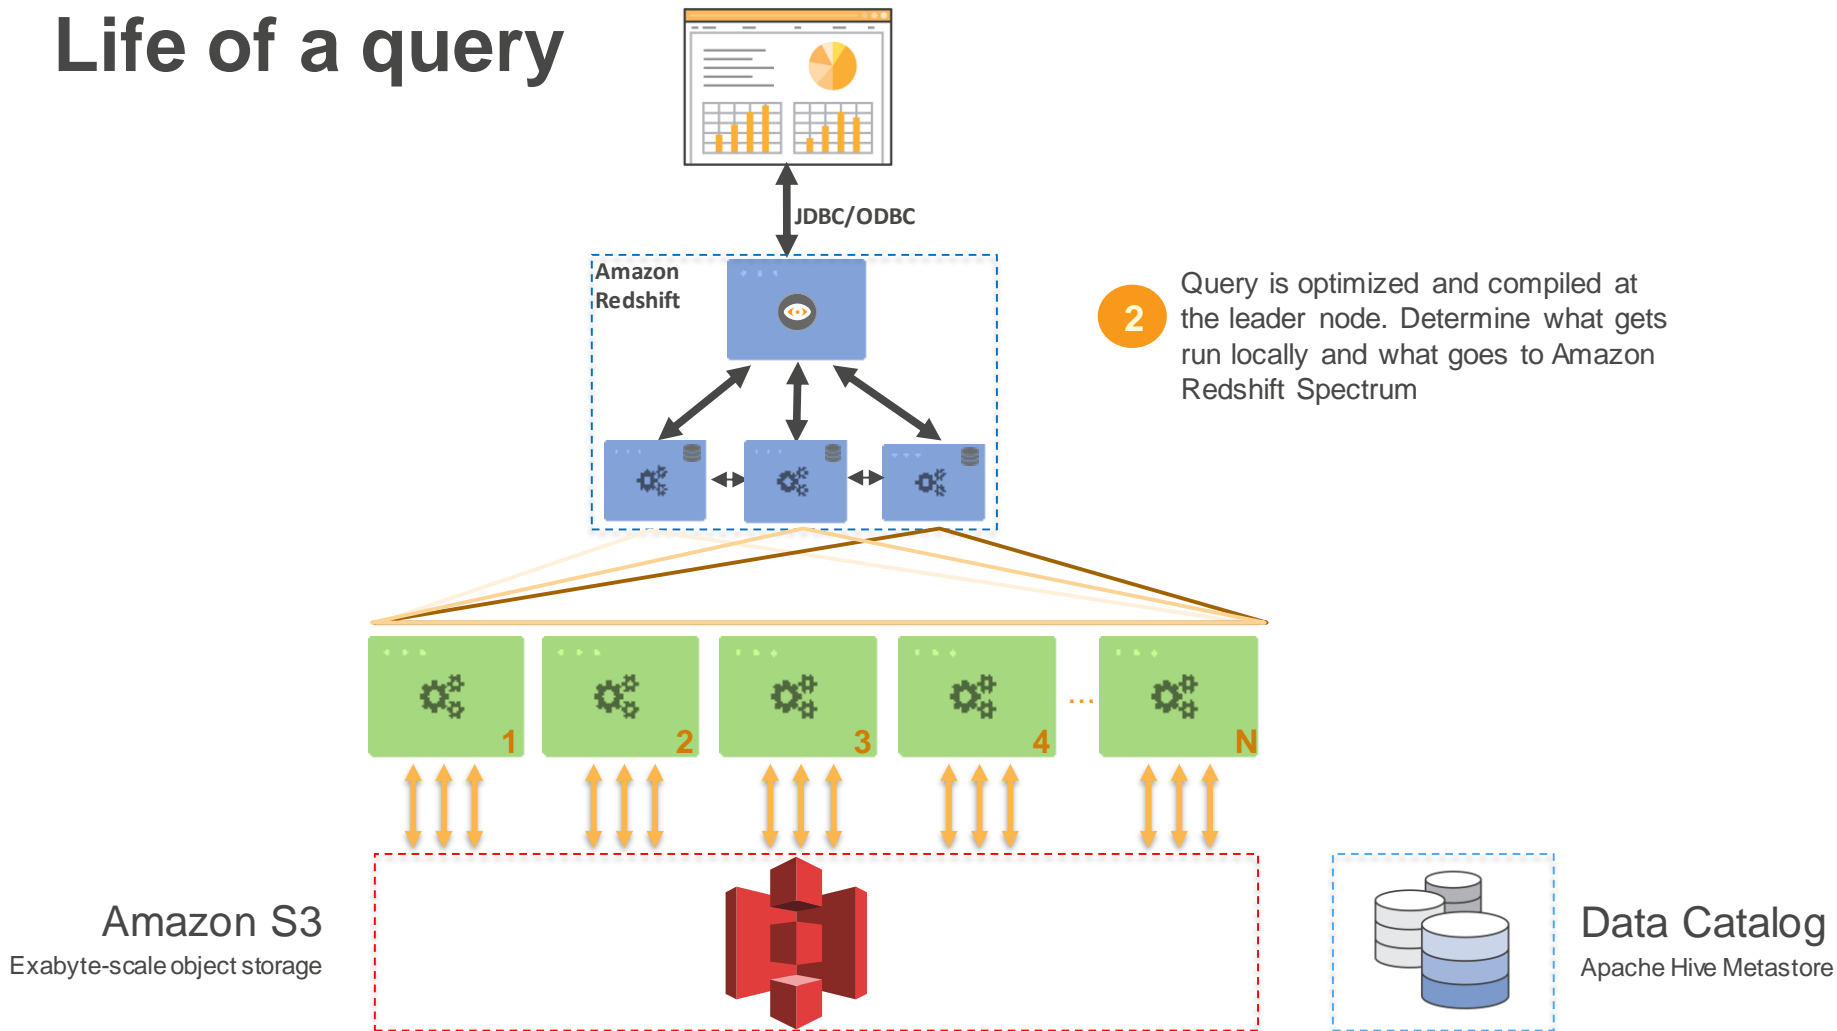


Full Amazon Redshift SQL support

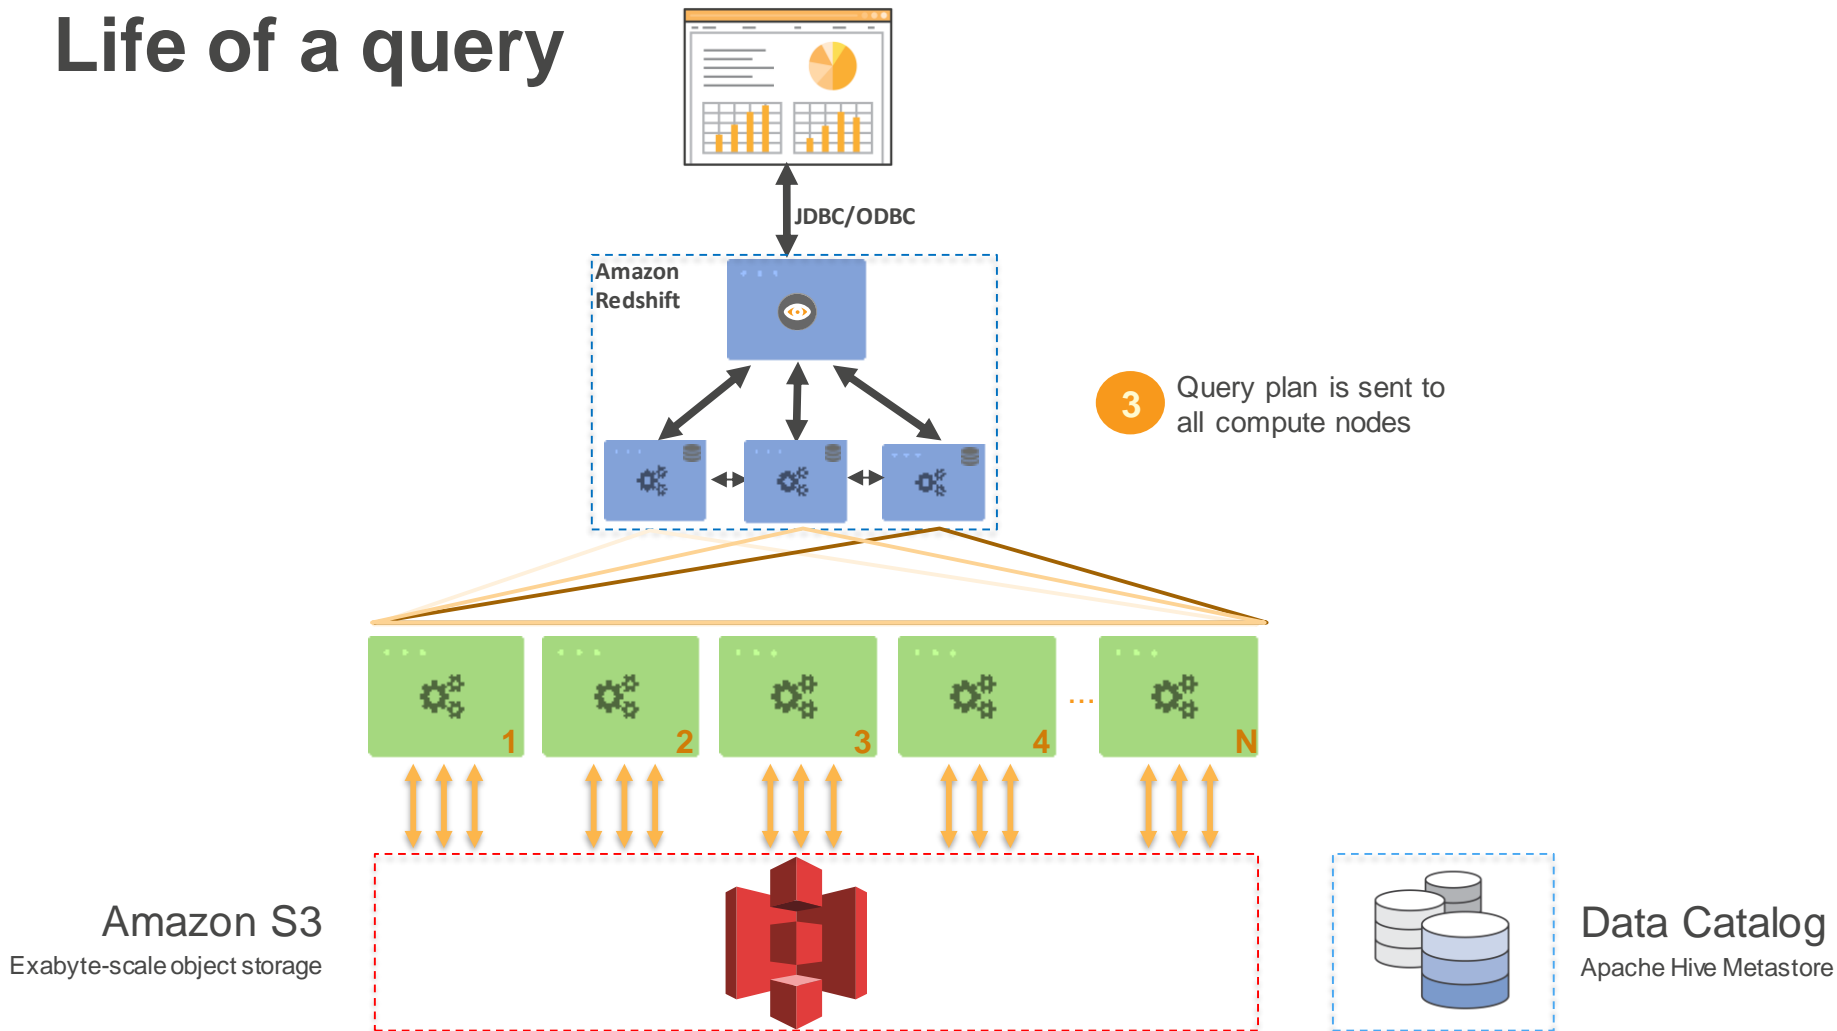
Life of a query



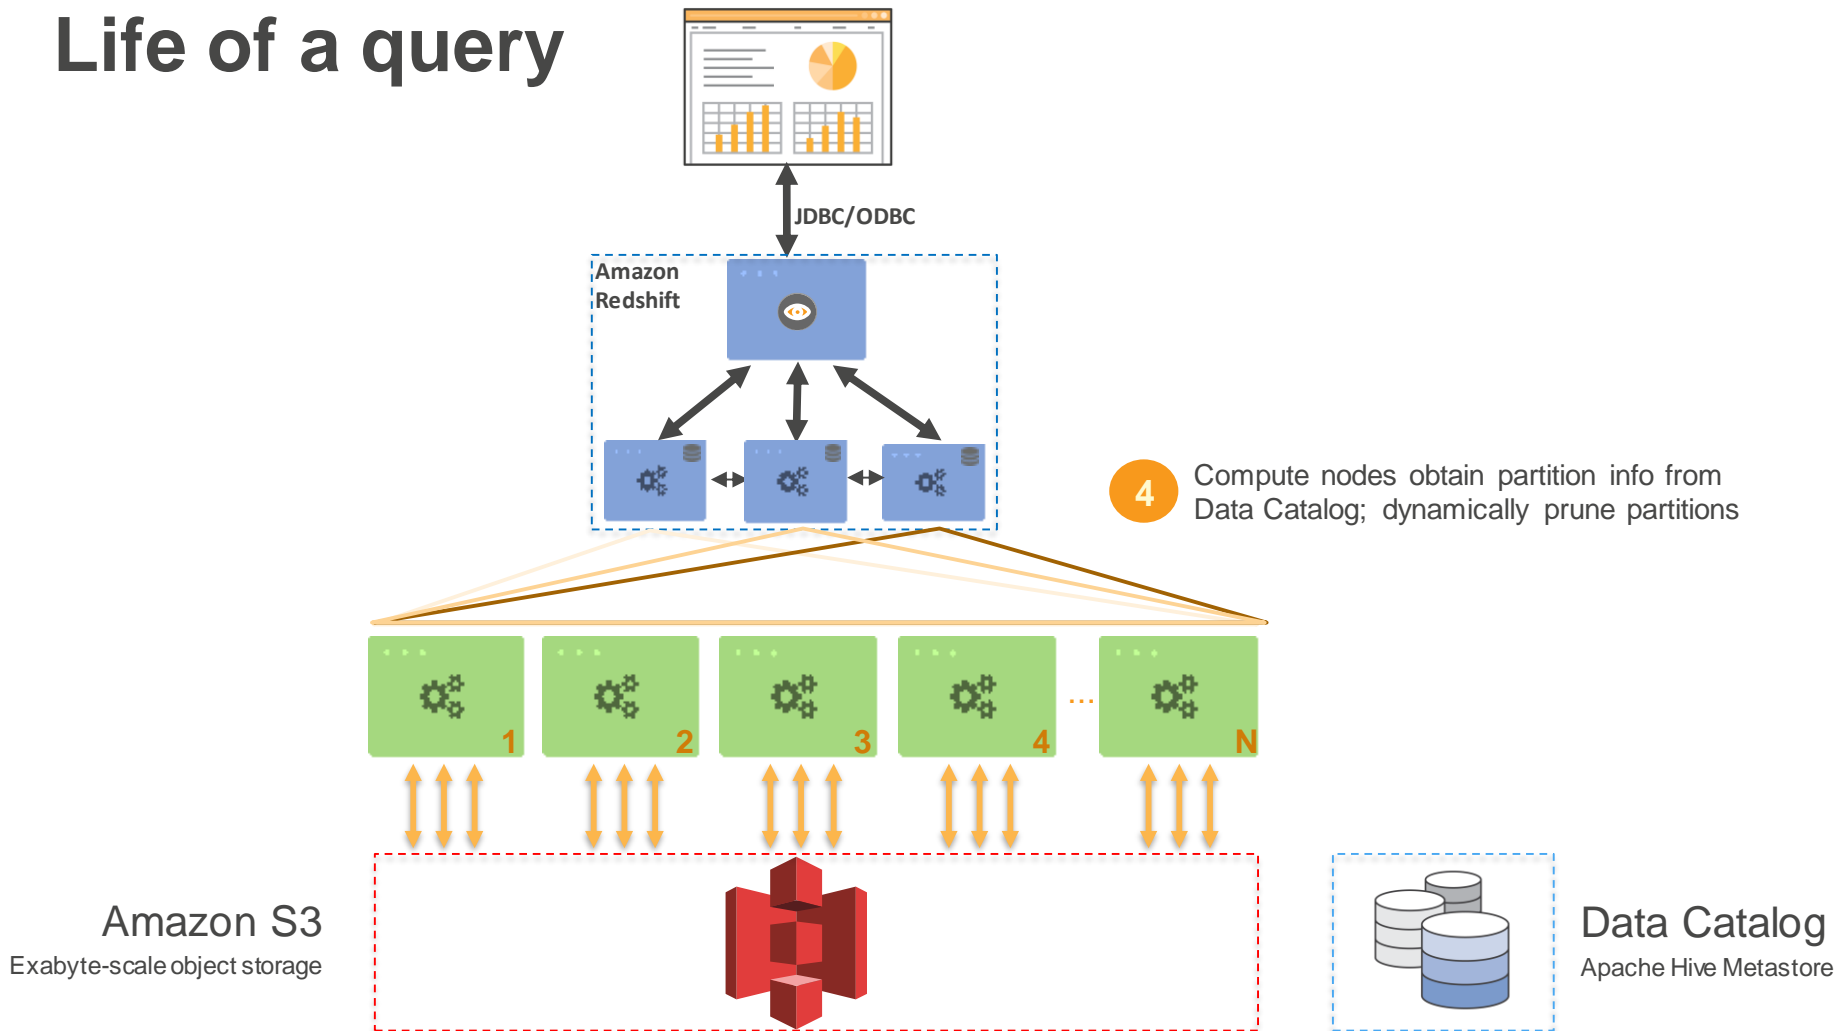
Life of a query



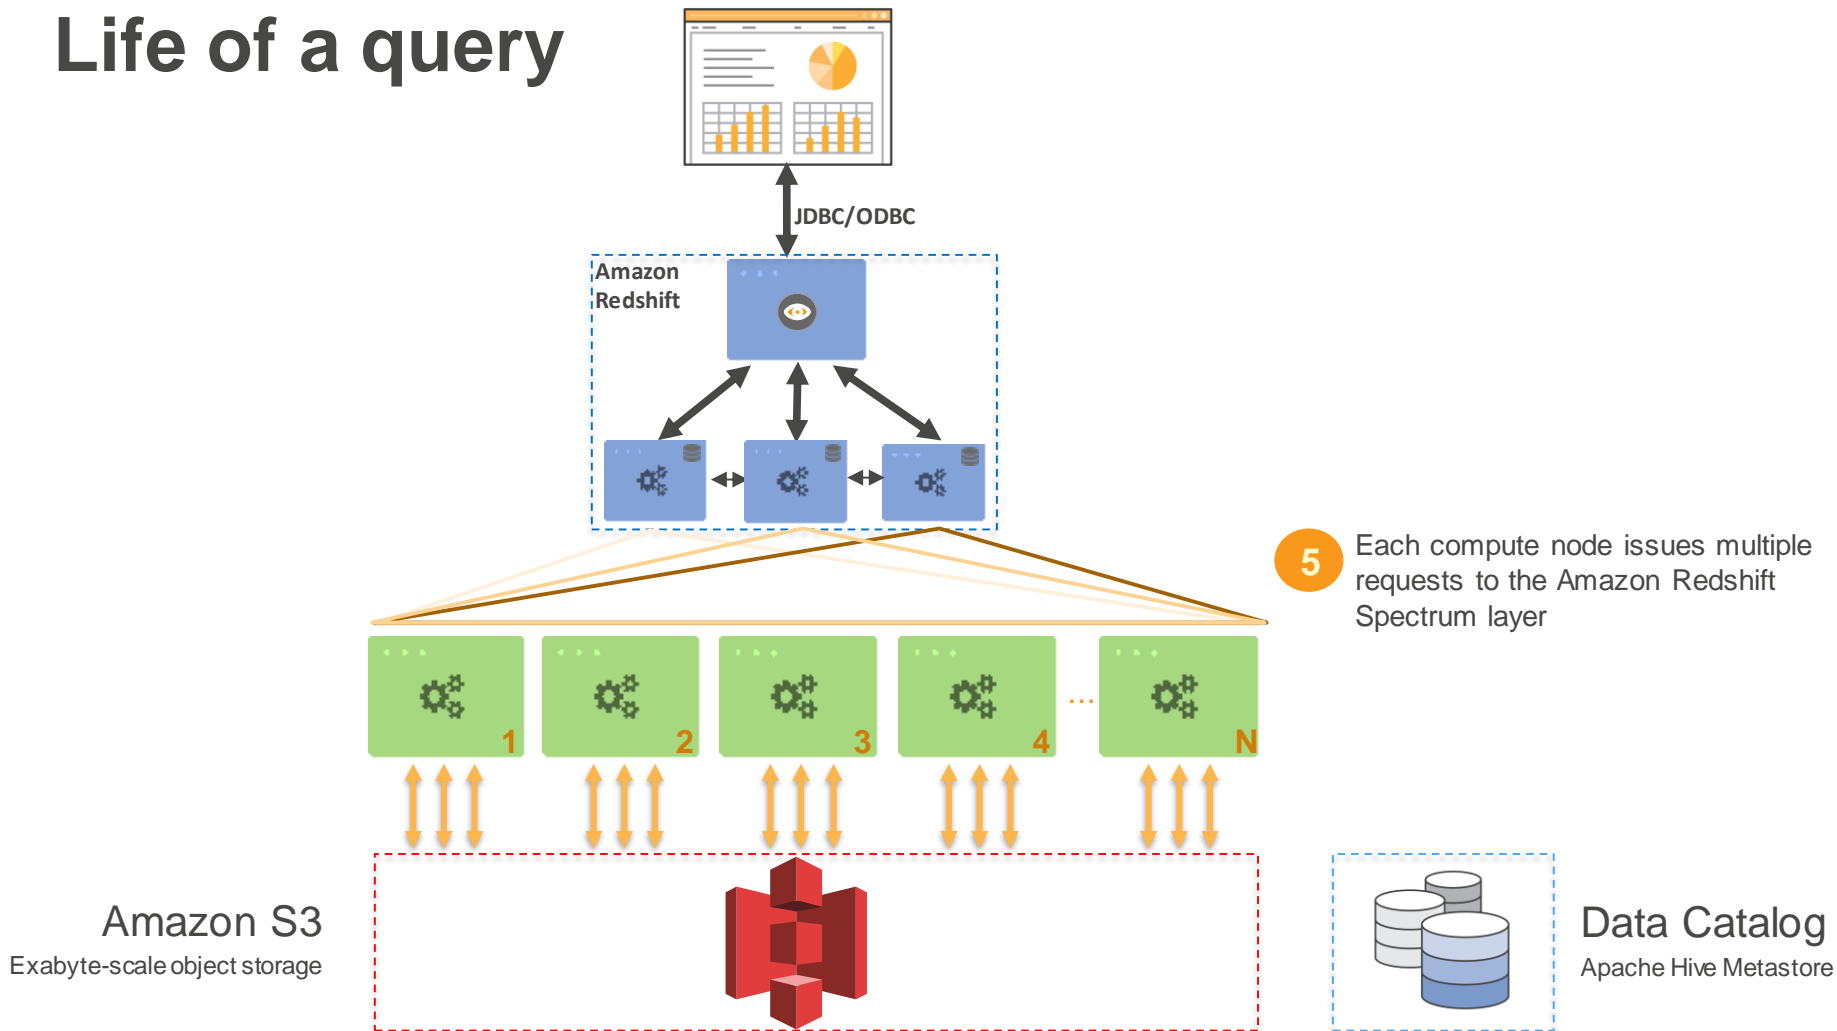
Life of a query



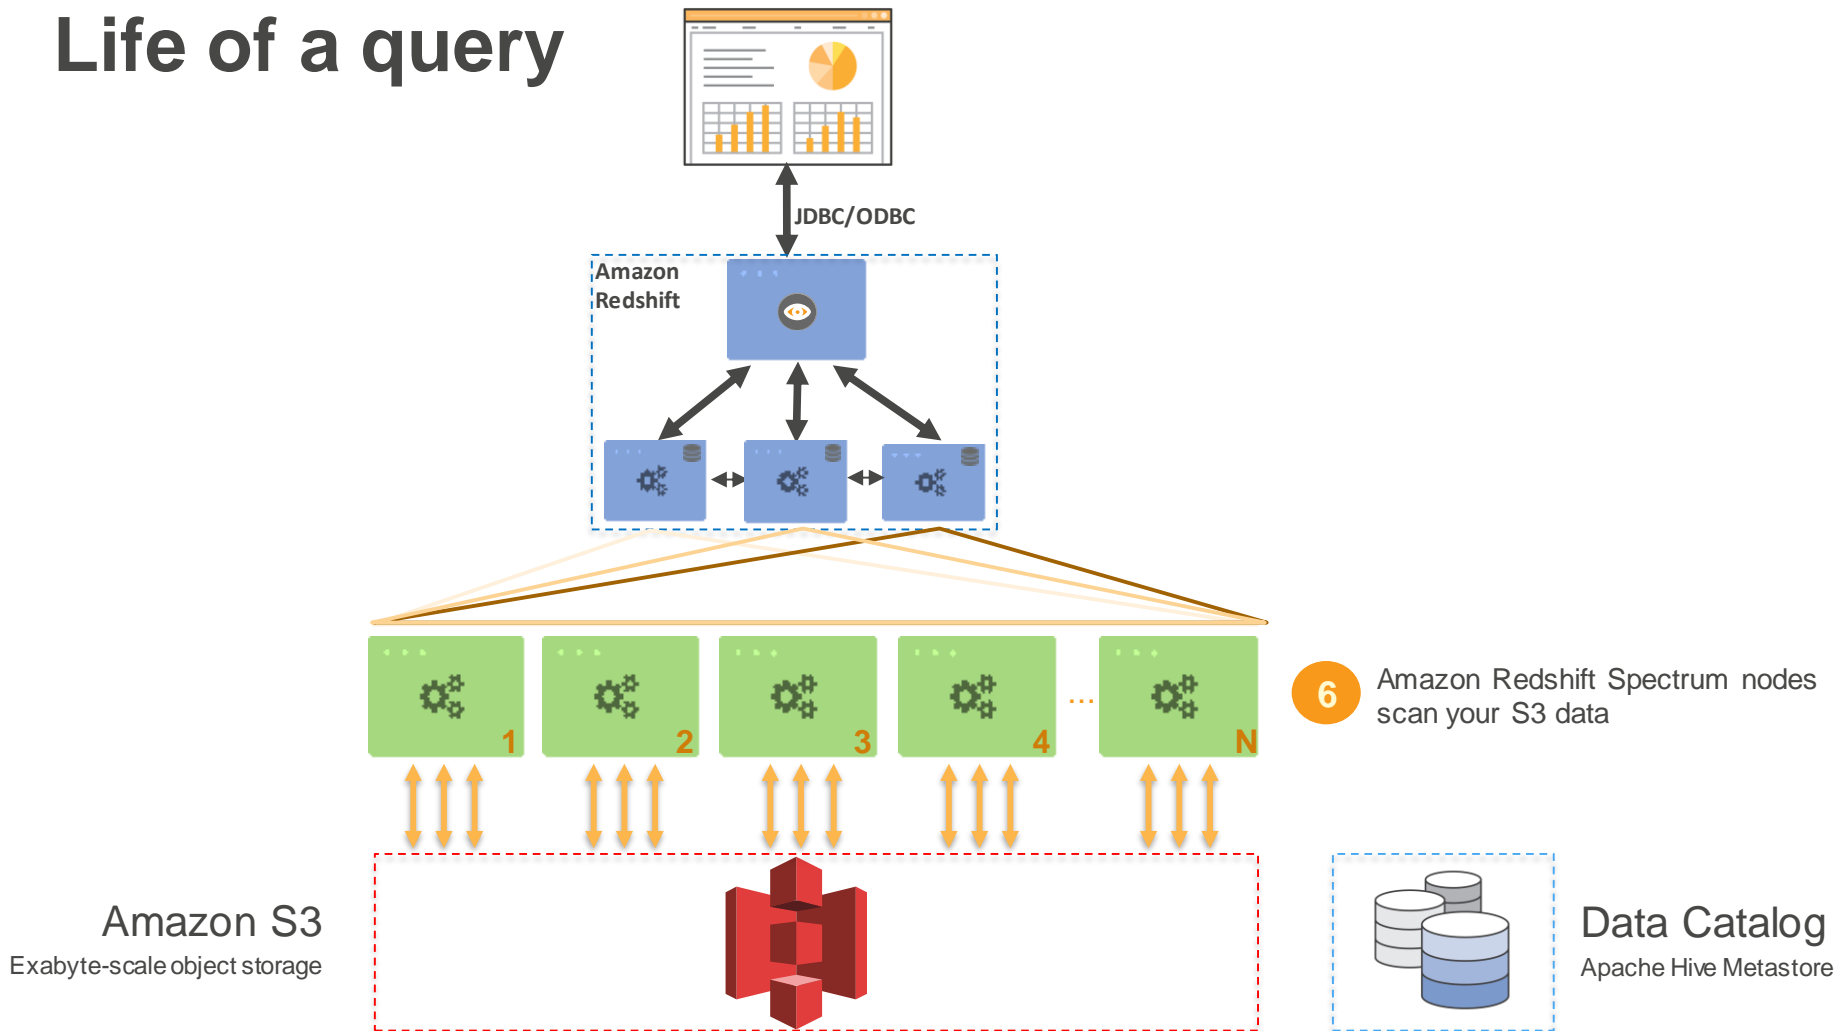
Life of a query



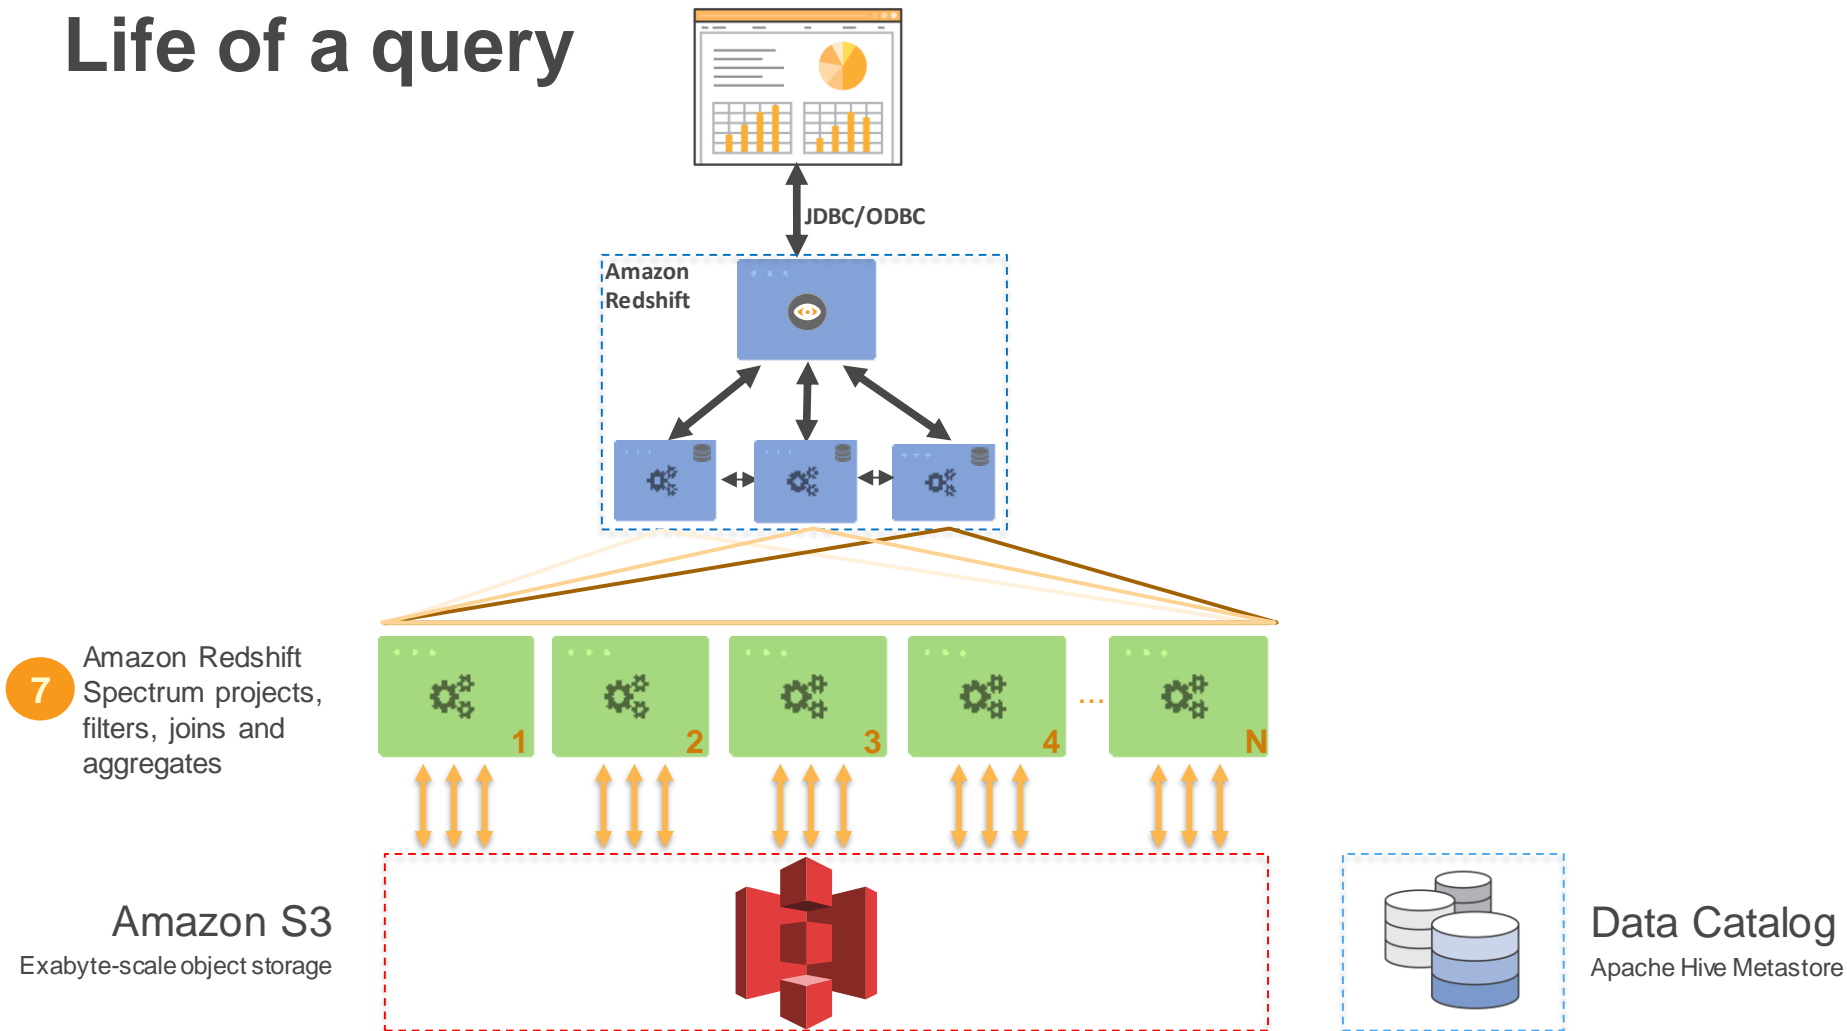
Life of a query



Life of a query




Life of a query



Life of a query



JDBC/ODBC



8

Final aggregations and joins with local Amazon Redshift tables done in-cluster



Amazon S3

Exabyte-scale object storage

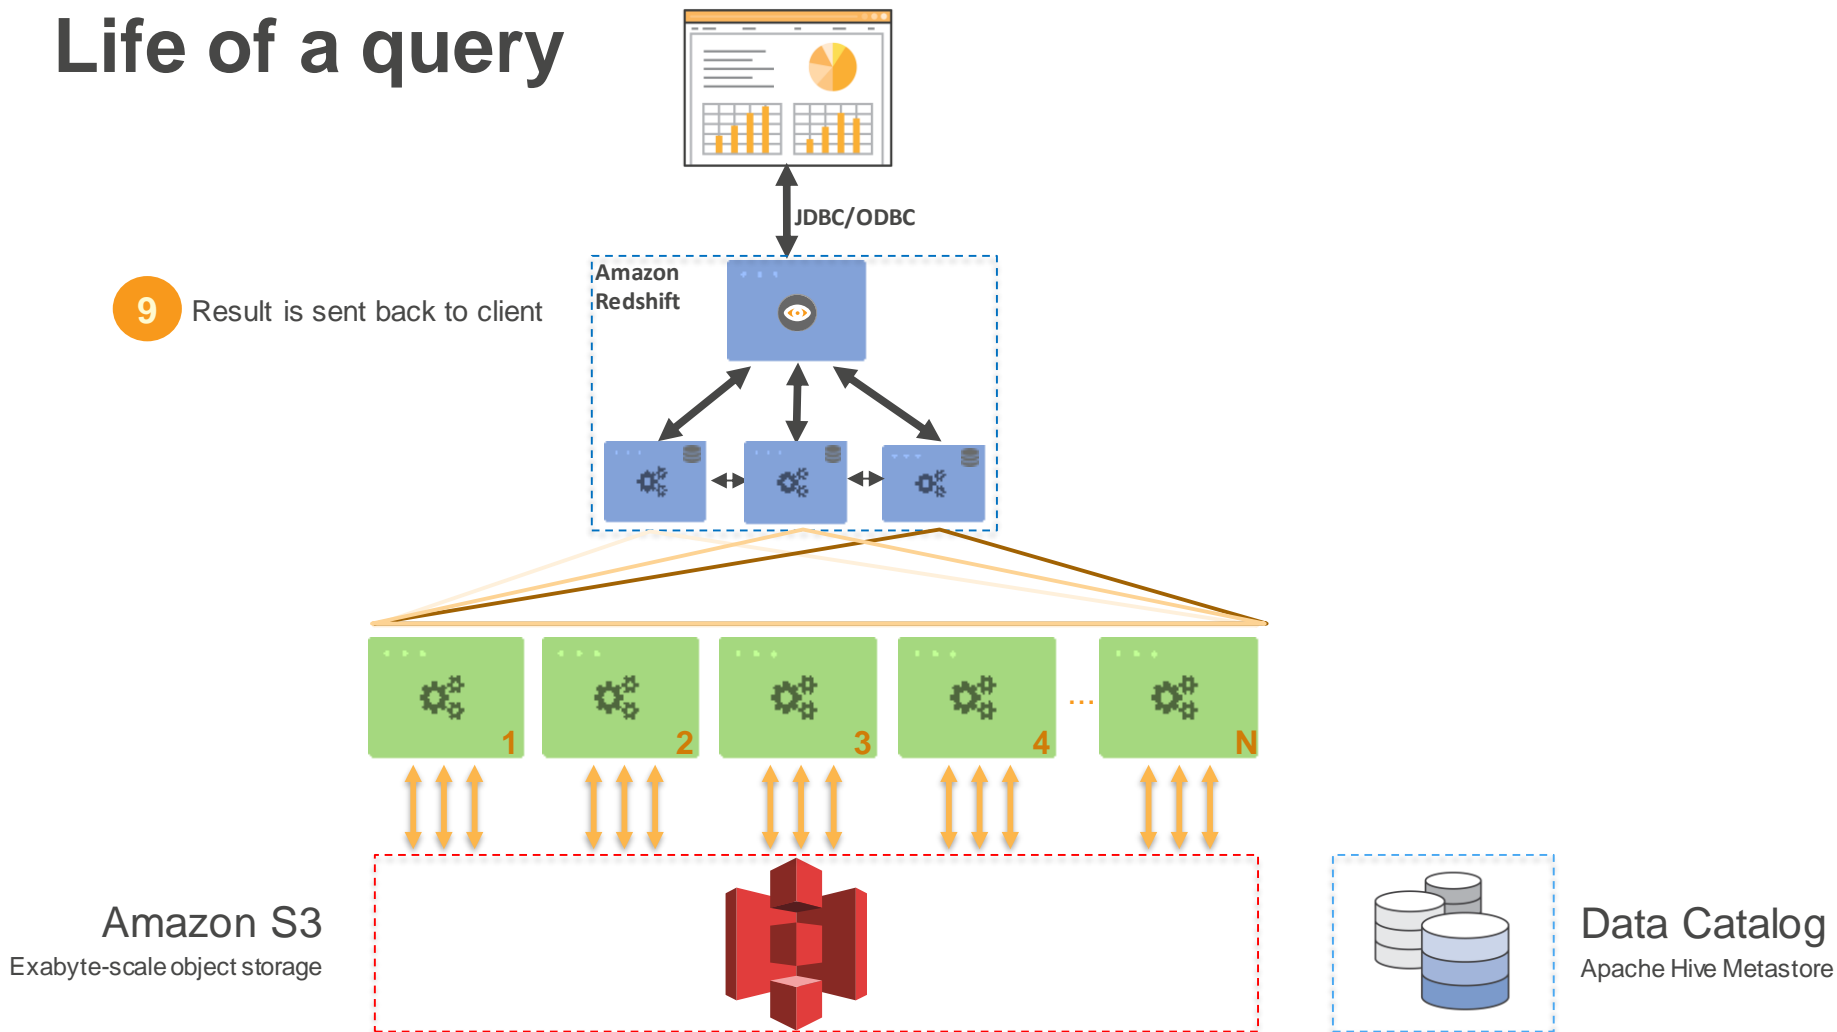


Data Catalog

Apache Hive Metastore

Life of a query

9 Result is sent back to client



Demo:

Running an analytic query over an exabyte in S3

Lets build an analytic query - #1

An author is releasing the 8th book in her popular series. How many should we order for Seattle? What were prior first few day sales?

Lets get the prior books she's written.

1 Table

2 Filters

```
SELECT
    P.ASIN,
    P.TITLE
FROM
    products P
WHERE
    P.TITLE LIKE '%POTTER%' AND
    P.AUTHOR = 'J. K. Rowling'
```

Lets build an analytic query - #2

An author is releasing the 8th book in her popular series. How many should we order for Seattle? What were prior first few day sales?

Lets compute the sales of the prior books she's written in this series and return the top 20 values

2 Tables (1 S3, 1 local)

2 Filters

1 Join

2 Group By columns

1 Order By

1 Limit

1 Aggregation

```
SELECT
    P.ASIN,
    P.TITLE,
    SUM(D.QUANTITY * D.OUR_PRICE) AS SALES_sum
FROM
    s3.d_customer_order_item_details D,
    products P
WHERE
    D.ASIN = P.ASIN AND
    P.TITLE LIKE '%Potter%' AND
    P.AUTHOR = 'J. K. Rowling' AND
GROUP BY P.ASIN, P.TITLE
ORDER BY SALES_sum DESC
LIMIT 20;
```

Lets build an analytic query - #3

An author is releasing the 8th book in her popular series. How many should we order for Seattle? What were prior first few day sales?

Lets compute the sales of the prior books she's written in this series and return the top 20 values, just for the first three days of sales of first editions

3 Tables (1 S3, 2 local)

5 Filters

2 Joins

3 Group By columns

1 Order By

1 Limit

1 Aggregation

1 Function

2 Casts

```
SELECT
    P.ASIN,
    P.TITLE,
    P.RELEASE_DATE,
    SUM(D.QUANTITY * D.OUR_PRICE) AS SALES_sum
FROM
    s3.d_customer_order_item_details D,
    asin_attributes A,
    products P
WHERE
    D.ASIN = P.ASIN AND
    P.ASIN = A.ASIN AND
    A.EDITION LIKE '%FIRST%' AND
    P.TITLE LIKE '%Potter%' AND
    P.AUTHOR = 'J. K. Rowling' AND
    D.ORDER_DAY :: DATE >= P.RELEASE_DATE AND
    D.ORDER_DAY :: DATE < dateadd(day, 3, P.RELEASE_DATE)
GROUP BY P.ASIN, P.TITLE, P.RELEASE_DATE
ORDER BY SALES_sum DESC
LIMIT 20;
```

Lets build an analytic query - #4

An author is releasing the 8th book in her popular series. How many should we order for Seattle? What were prior first few day sales?

Lets compute the sales of the prior books she's written in this series and return the top 20 values, just for the first three days of sales of first editions in the city of Seattle, WA, USA

4 Tables (1 S3, 3 local)

8 Filters

3 Joins

4 Group By columns

1 Order By

1 Limit

1 Aggregation

1 Function

2 Casts

```
SELECT
    P.ASIN,
    P.TITLE,
    R.POSTAL_CODE,
    P.RELEASE_DATE,
    SUM(D.QUANTITY * D.OUR_PRICE) AS SALES_sum
FROM
    s3.d_customer_order_item_details D,
    asin_attributes A,
    products P,
    regions R
WHERE
    D.ASIN = P.ASIN AND
    P.ASIN = A.ASIN AND
    D.REGION_ID = R.REGION_ID AND
    A.EDITION LIKE '%FIRST%' AND
    P.TITLE LIKE '%Potter%' AND
    P.AUTHOR = 'J. K. Rowling' AND
    R.COUNTRY_CODE = 'US' AND
    R.CITY = 'Seattle' AND
    R.STATE = 'WA' AND
    D.ORDER_DAY :: DATE >= P.RELEASE_DATE AND
    D.ORDER_DAY :: DATE < dateadd(day, 3, P.RELEASE_DATE)
GROUP BY P.ASIN, P.TITLE, R.POSTAL_CODE, P.RELEASE_DATE
ORDER BY SALES_sum DESC
LIMIT 20;
```

Now let's run that query over an exabyte of data in S3

```
demo=# SELECT
demo-#   P.ASIN,
demo-#   P.TITLE,
demo-#   R.POSTAL_CODE,
demo-#   P.RELEASE_DATE,
demo-#   SUM(D.QUANTITY * D.OUR_PRICE) AS SALES_sum
demo-# FROM s3.d_customer_order_item_details D, asin_attributes A, products P, regions r
demo-# WHERE D.ASIN = P.ASIN AND
demo-#        P.ASIN = A.ASIN AND
demo-#        D.REGION_ID = R.REGION_ID AND
demo-#        A.EDITION LIKE '%FIRST%' AND
demo-#        P.TITLE LIKE '%Potter%' AND
demo-#        P.AUTHOR = 'J. K. Rowling' AND
demo-#        R.COUNTRY_CODE = 'US' AND
demo-#        R.CITY = 'Seattle' AND
demo-#        R.STATE = 'WA' AND
demo-#        D.ORDER_DAY :: DATE >= P.RELEASE_DATE AND
demo-#        D.ORDER_DAY :: DATE < dateadd(day, 3, P.RELEASE_DATE)
demo-# GROUP BY P.ASIN, P.TITLE, R.POSTAL_CODE, P.RELEASE_DATE
demo-# ORDER BY sales_sum DESC
demo-# LIMIT 20;[]
```

Roughly 140 TB of customer item order detail records for each day over past 20 years.

190 million files across 15,000 partitions in S3. One partition per day for USA and rest of world.

Need a **billion-fold** reduction in data processed.

Running this query using a 1000 node Hive cluster would take over 5 years.*

- Compression5X
- Columnar file format.....10X
- Scanning with 2500 nodes.....2500X
- Static partition elimination.....2X
- Dynamic partition elimination.....350X
- Redshift's query optimizer.....40X

Total reduction.....**3.5B X**

* Estimated using 20 node Hive cluster & 1.4TB, assume linear

* Query used a 20 node DC1.8XLarge Amazon Redshift cluster

* Not actual sales data - generated for this demo based on data format used by Amazon Retail.

Is Amazon Redshift Spectrum useful if I don't have an exabyte?

Your data will get bigger

On average, data warehousing volumes grow 10x every 5 years

The average Amazon Redshift customer doubles data each year



Amazon Redshift Spectrum makes data analysis simpler

Access your data without ETL pipelines

Teams using Amazon EMR, Athena & Redshift can collaborate using the same data lake

Amazon Redshift Spectrum improves availability and concurrency

Run multiple Amazon Redshift clusters against common data

Isolate jobs with tight SLAs from ad hoc analysis

Architecture

Tuning

Integration

Spectrum

Echo
System


MAG

Summary

Redshift Partner Echo System

4 types of partners

- Load and transform your data with [Data Integration Partners](#)
- Analyze data and share insights across your organization with [Business Intelligence Partners](#)
- Architect and implement your analytics platform with [System Integration and Consulting Partners](#)
- Query, explore and model your data using tools and utilities from [Query and Data Modeling Partners](#)



aws.amazon.com/redshift/partners/

“Some” Amazon Redshift Customers



BEACHMINT



NOKIA



Manchester Airport Group

An AWS Redshift customer story





MAG – take-off with cloud and data

Stuart Hutson – Head of Data and BI

THE AVIATION PROFESSIONALS

MAG is a leading UK based airport company, which owns and operates Manchester, London Stansted, East Midlands and Bournemouth airports.

MAG is privately managed on behalf of its shareholders, the local authorities of Greater Manchester and Industry Funds Management (IFM). IFM is a highly experienced, long-term investor in airports and already has significant interests in ten airports across Australia and Europe.

48.5 MILLION passengers served per year.

Over **80 AIRLINES** serving **272 DESTINATIONS** direct.

£134.3 MILLION RETAIL INCOME per annum delivered via 200+ shops, bars and restaurants.


£125.7 MILLION CAR PARKS INCOME delivered via 96,000 parking spaces.

£623 MILLION property assets across all airports, 5.67m sq ft of commercial property.

£738.4 MILLION REVENUE +10.0% increase from last year.

£283.6 MILLION EBITDA growth of 17.2% in 2015.

£5.6 BILLION contribution to the UK economy from MAG airports.



OUR AIRPORTS...

MAG airports serve over 48.5 million people per annum from complementary catchment areas covering over 75% of the UK population.



c. 23m passengers per annum.
UK's 3rd largest airport.
70+ airlines & 200+ destinations.
2 runways with potential 62% capacity.
21.5m people within a 2 hour drive.




c. 0.7m passengers per annum.
Significant investment in new terminal increasing passenger capacity to 3m p.a.
Wealthy catchment area.
Large land holding – on-site business park.



c. 4.5m passengers per annum.
UK's largest freight airport after Heathrow – 310,000 tonnes p.a.
Located next to key road interchanges – four hours from virtually all UK commerce.



c. 23m passengers per annum.
UK's 4th largest airport.
150+ destinations.
1 runway with 50% spare capacity.
25m people within 2 hour drive.
Acquired February 2015.



Catchment area within 2 hours' drive of:



OUR CONNECTIVITY...

80+ airlines and over 270 direct destinations providing global connectivity.



AIR SERVICE DEVELOPMENT

MAG has a diverse carrier mix from global destinations with an excellent track record of incentivising passenger growth.

MAG has exceeded expectations with industry leading rates of passenger growth. Importantly for passengers, by forging strong commercial partnerships with airlines, our airports have been able to increase choice and convenience and make a stronger contribution to economy growth.

CARGO SERVICE DEVELOPMENT

MAG's Cargo produces an annual income of £20.2 million and holds 26% of the UK freight market share.

East Midlands is the UK's largest dedicated freight hub handling 310,000 tonnes of freight per annum. Stansted handles 233,000 tonnes of freight per annum and is a key gateway to London and the South of England.

OUR DEVELOPMENTS...

Manchester Transformation Programme and London Stansted Transformation Programme are developments that all aim to drive improved customer service.

MANCHESTER TRANSFORMATION PROGRAMME

With investment of £1 billion, Manchester will become one of the most modern and customer focused airports in Europe demonstrating the importance of Manchester as a global gateway.

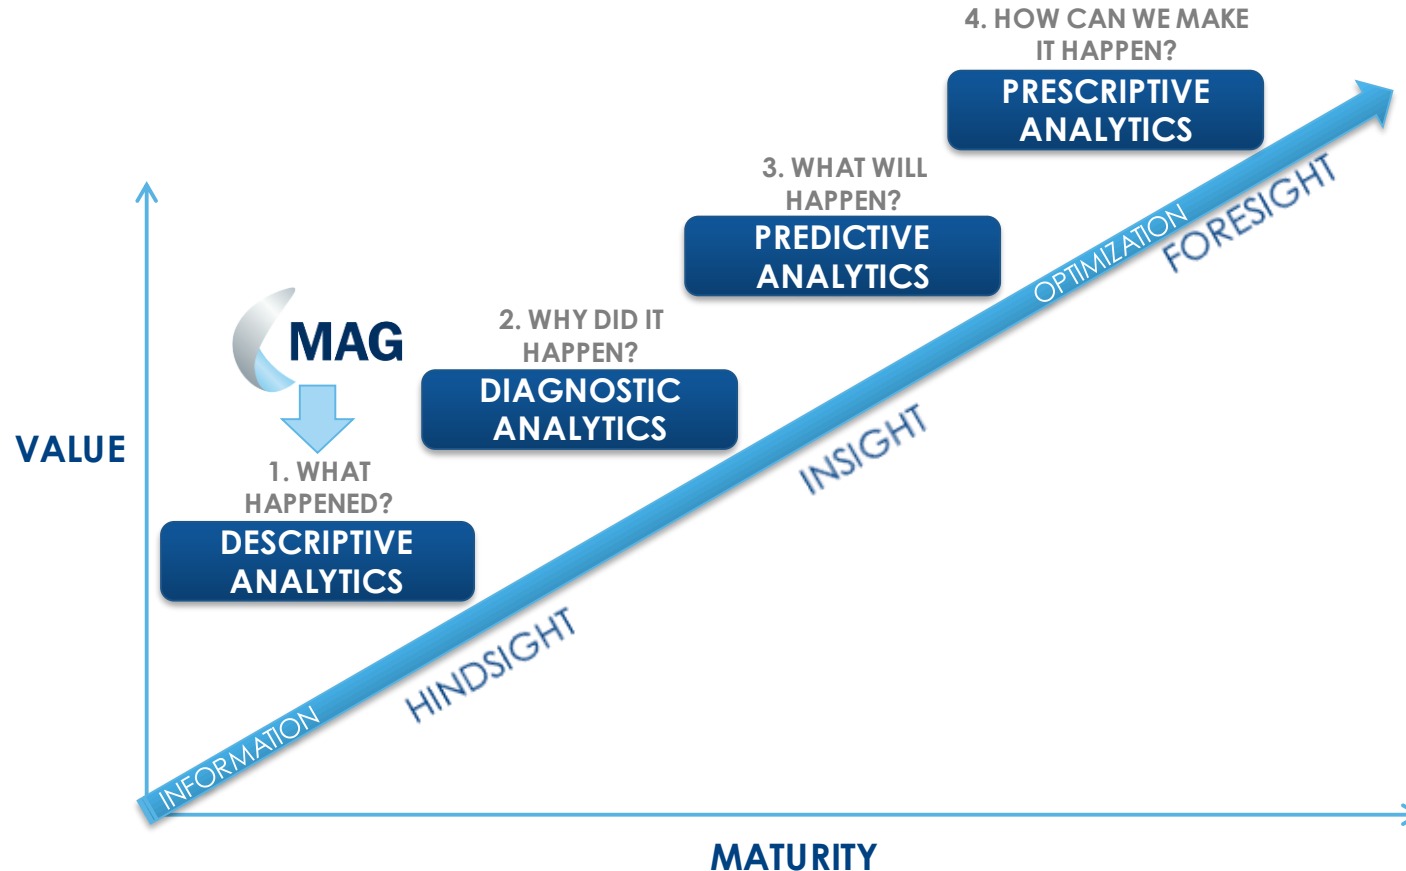


LONDON STANSTED TRANSFORMATION PROGRAMME

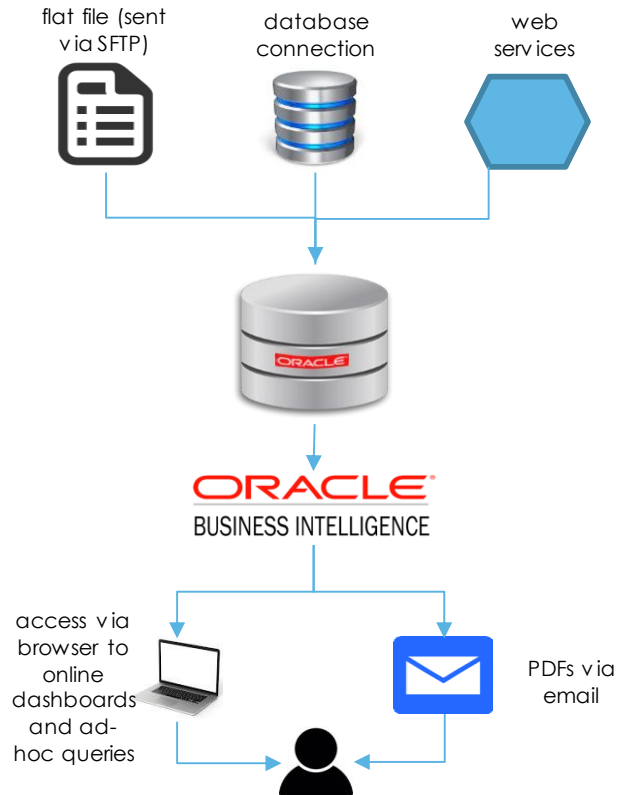
The £80 million terminal transformation project at London Stansted will transform the passenger experience and boost commercial yields.



MAG'S CURRENT BUSINESS INTELLIGENCE MATURITY



MAG'S LEGACY ARCHITECTURE - CHALLENGES...



Business

- Multiple version of the truth
- Unclear of operational problems
- People are overloaded and with data and data led questions
- Analysts not able to do analyst job due to lack of data and tools
- Data processing issues - late reports, missing data
- Data accessible in silos - no real cross-functional analysis

Technical

- Database @ 95% capacity on physical kit that can not be scaled.
- Dashboards are slow to run.
- Constant optimisation and maintenance of database.
- Limited concurrent connections for queries.
- Lack of self-serve – centralised BI model.
- No direct connection to database – business wants to expand into using R and Python etc.
- All data in batch with no possibility of streaming

VALUE OF BI STRATEGY IN MAG...

Monetise Data

- **Monetise data and technology across our omni-channels:** *MAG's BI Strategy must be bold, it should be aiming for how we monetises our data and technology across our omni-channel business by improving the customer experience.*

Democratise Data

- **Democratise data across the Enterprise:** *Our data needs to be pervasive across the organisation. The decisions of the organisation should be made on clear information presented to the business at the right time to enable MAG to make the right decisions.*

Data DNA

- **Create a data DNA:** *Build a culture around data and analytical thinking across the organisation by embedding analytics and data across MAGs business processes and decision making.*

PHASE 1 - IMPLEMENT SELF-SERVE RETAIL BI SOLUTION

- 50+ PARTNERS GENERATING OVER £130M REVENUE...

Build Data and BI Foundation solution

- To create an extensible and flexible data solution for MAG comprising of:
 - **Extended Data Warehouse.**
 - Scalable and elastic compute.
 - Deal with seasonality spikes of passenger travel.
 - **Real-time streaming.**
 - Enable MAG to become a real-time business across their customer journey.
 - **Cloud environment:**
 - Secured.
 - Resilient.
 - Repeatable build.
- **Enable MAG to quickly experiment at low cost and minimal risk.**
 - MAG wants to trial new technologies, especially open-source.
- **Create an architecture that can evolve over time to meet MAG's new challenges.**
 - Benefits delivered early and continuously.
 - No need for MAG to invest in a large, front-loaded EDW programme.



EXAMPLE OF MAG'S DESIGN PRINCIPLES TO SOLVE THE PROBLEMS...

- **Evolutionary architecture**
- **Infrastructure as Code**
- **Protecting our data**
- Assume for failure
- Data quality is a priority
- Embrace open source for experimentation
- SaaS -> PaaS -> IaaS
- Serverless computing
- Etc.



MAG – OUR 6 MONTH JOURNEY...

From		To
Single instance database.	→	Scale-able Data Warehouse.
Daily sales rung in at store level.	→	Over 90% of all sales automatically ingested at product level.
Car parking - flat files ingested in batch.	→	Ingest and interrogate streaming data directly: <ul style="list-style-type: none">• Car park data is being added via Kinesis
Access to database limited to reporting tool.	→	Authorised users can use visualisation and data science tools (e.g. R and Python) of their choice for self-serve analytics
No database writeback for end-users.	→	Sandboxes in Redshift for end user experimentation.

MAG – NEXT 6-12 MONTHS...

- Moving to near-time streaming into Redshift for:
 - Terminal Operations
 - Security Services
 - Car Park Management
- Streaming semi-structured data into Redshift
 - Trialling IoT data streaming
 - Passenger analysis
- Trial AWS Glue and AWS Redshift Spectrum
 - Automated profile and catalogue of data across the enterprise
 - Continuous integration of data into our data warehouse

The background is a full-page image with a red color overlay. It depicts a calm body of water, possibly a lake or a wide river, with a hazy, mountainous shoreline in the distance. In the foreground, on the right side, a group of four people are standing on a cluster of dark, jagged rocks that protrude from the water. The people are silhouetted against the lighter background. The overall mood is serene and contemplative.

Who Are Crimson Macaw

Driving customer success by unlocking the value of data.

Competency focused consultancy



Architecture
and
Data Strategy



DevOps



Data Engineering



Enterprise
Data
Solutions



Data
Science

1. Plan

2. Build

3. Action

www.crimsonmacaw.com

Our partners ...



3 AWS Big Data Speciality

3 AWS Certified Solutions Architect Associate

2 AWS Certified Developer Associate

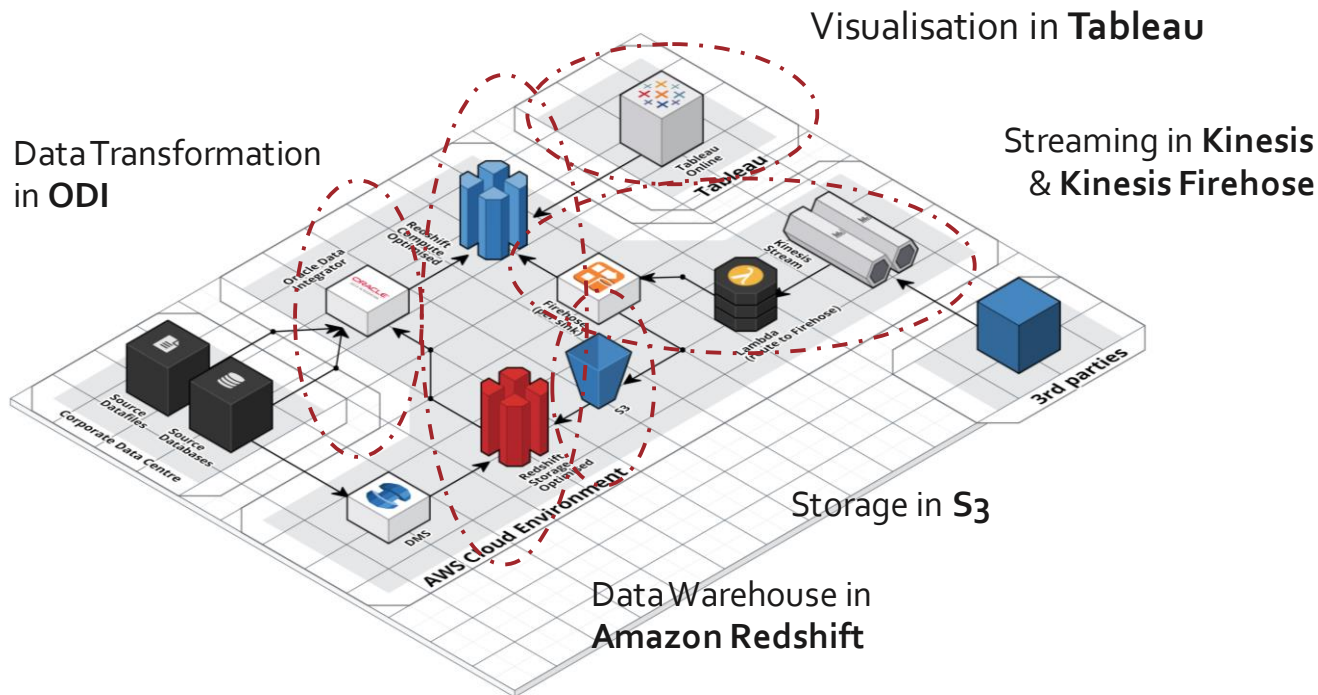
2 AWS Certified SysOps Administrator Associate



Building a solution

... without too many twists and turns.

Key architectural components used



Cloud Architecture + Data Architecture = Solution

How do you match the pace of infrastructure build in the cloud with understanding the data & BI requirements?

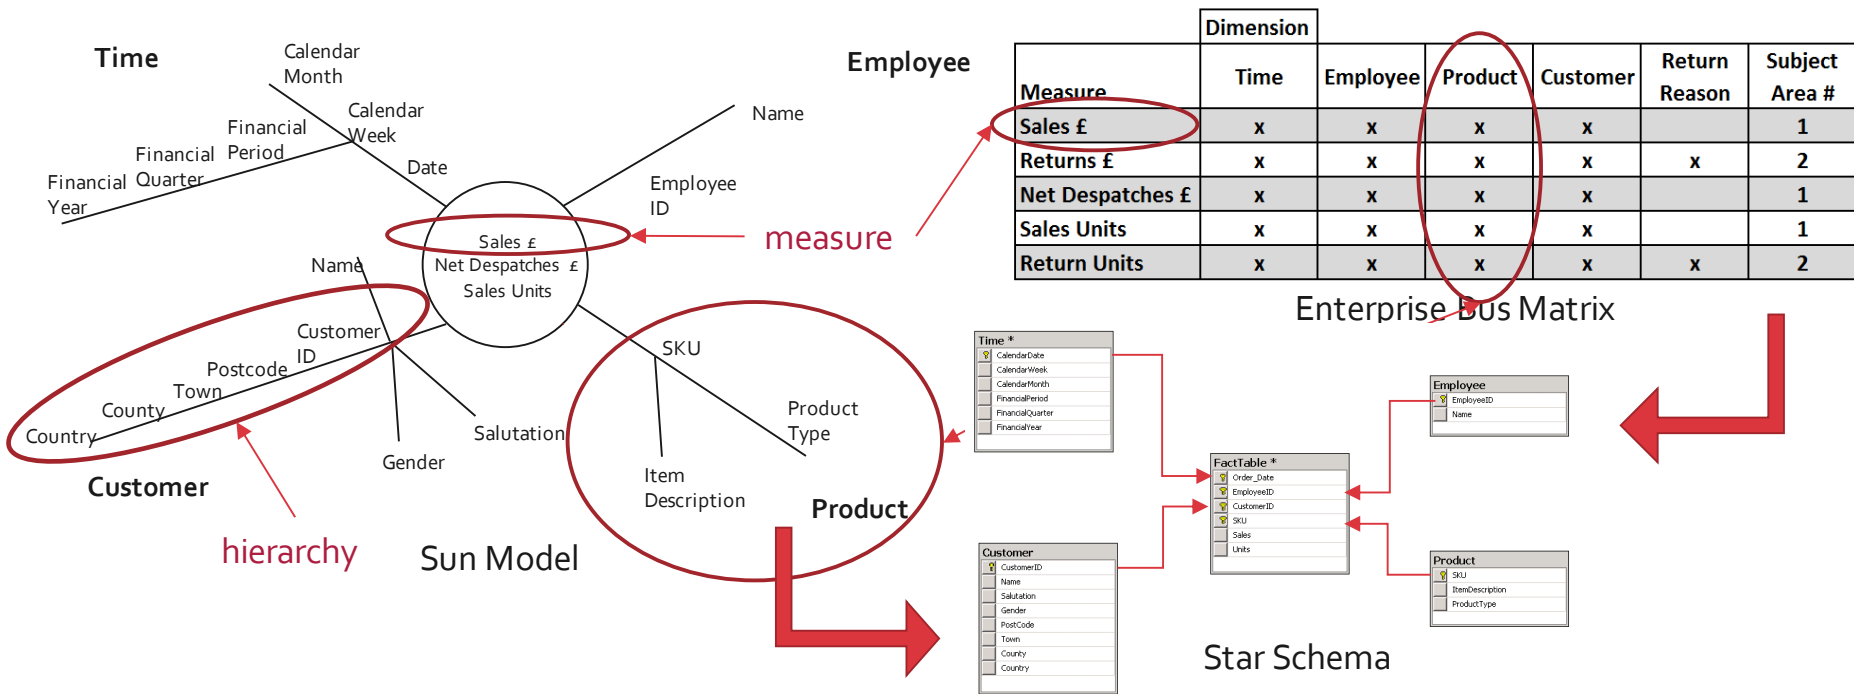
Deliver value quickly vs conformed dimensions?

- A horizontal analytical 'slice' across the estate.
- Understand conformed dimensions.
- Vertical slice of a business domain.
- Reduced refactoring due to the prior horizontal analysis.

Understand how the business will consume and use the data?

- Produce artefacts that are:
 - Shared by stakeholders and the delivery team.
 - Understandable by all parties.
 - Highly visual, allow complex information to be absorbed - sun modelling.

Sun modelling vs Enterprise Bus Matrix




Building the infrastructure (as code)



- Why use infrastructure as code?
 - Repeatability.
 - Consistency.
 - Versioned.
 - Code reviews.
 - Speed of delivery.
- Technology Used:
 - CloudFormation in YAML format with custom YAML Tags.
 - Lambda Functions for Custom Resource Types.
 - Bespoke deployment utility.
 - Puppet Standalone in Cloud Init for EC2.
- Why this approach?
 - Enforced Tagging Policy with propagated tags.
 - Custom YAML Tags act as a precompiler for CloudFormation.
 - Not all resources types were available, e.g. DMS.
 - Redshift IAM Roles and Tags – both now available out of the box!.

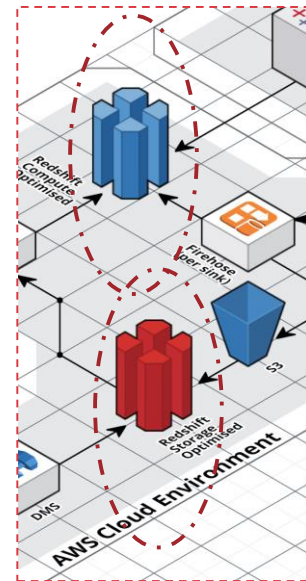
Security overview

- Three independent AWS accounts
 - Dev – for development of data processes.
 - Prod – target deployment.
 - Sec – sink for data generated by Config Service and CloudTrail to S3 buckets.
- Encryption
 - KMS Encryption keys used throughout
 - Enforced SSL connections to Redshift
 - S3 – enforced write encryption (by policy).
- Audit and compliance documentation
 - AWS Artifacts.




Redshift topology

- Storage Optimised (red)
 - Optimised for storing larger volumes of data (source).
 - Ingestion point for newly arriving data.
 - Transformation layer (large number of work tables).
 - VPC - private subnet.
- Compute Optimised layer (blue)
 - Transformed data.
 - Near real-time operational data.
 - Present dimensional layer.
 - VPC – public subnet (whitelisted access).



What about Streaming?

- Setup Kinesis Streams to allow 3rd parties to send data.
- Enabled cross account access with an assumed role.
- Used Lambda to route mixed data to multiple Firehose Streams.
- Firehose Streams sink data to S3 and/or Redshift Compute (blue).




Observations

A person wearing a black leather jacket and goggles is seen from behind, holding a vintage radio receiver. The radio has a circular face with various knobs, dials, and a speaker grille. The person is standing on a metal lattice structure, likely a radio tower, with a bright, hazy sky in the background. The entire image is overlaid with a red tint.

ODI and Redshift



- Problem: ODI initiated Redshift tasks not completing.




Solution: Increase Array Fetch Size in ODI

- Problem: No native knowledge modules in ODI for Redshift.
 - Solution: Customised existing generic SQL knowledge modules for Redshift.
 - Evaluating 3rd party solution Knowledge Module.

Tableau and Redshift



- How does Tableau Online connect to Redshift?
 - JDBC via SSL.
 - Whitelisted to Redshift.
 - Tableau available in multiple regions (US, Ireland).
- Enable Redshift constraints:
 - Foreign Key and Primary Key and Unique constraints – ensure they are created in Redshift (even though they are not enforced).
- Enable Tableau “Assume Referential Integrity”
 - in Tableau workbooks (if you have it!).
- Queries in Tableau:
 - Executed via Redshift cursor – minimise IO.
 - Current activity: `stv_active_cursors`.
 - For recent activity (two - five days): `stl_query` and `stl_utility_text`.



```
dev_munsoor_negyal    fetch 10000 in "SQL_CUR0x7fdcdb144800"
```

Tableau – getting back to SQL

```
select
    a.query,a.querytxt as cursor_sql
    ,b.sequence,b.text as raw_sql, b.starttime
from
    stl_query a inner join stl_utilitytext b
on a.pid = b.pid and a.xid = b.xid
where
    database = '<DBName>'
and a.starttime >= dateadd(day, -1, current_date)
order by
    a.xid, b.starttime, b.sequence asc;
```

query	cursor_sql	sequence	raw_sql	starttime
1009328 fetch 1000 in "SQL_CUR3";		0	close "SQL_CUR3"	2017-06-
1009329 Undoing 1 transactions on table		0	BEGIN;	2017-06-
1009329 Undoing 1 transactions on table		0	ROLLBACK;	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		0	begin;	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		0	declare "SQL_CUR6" cursor with hold for SELECT TOP 1000 1 AS "number of records", "fact_item_sales"."basket_properties_id" AS "basket_prope...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		1	card_id", "fact_item_sales"."concession_rate" AS "concession_rate", "fact_item_sales"."final_destination_airport_id" AS "final_destination_airport_id",...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		2	id" AS "flight_actual_departure_date_id", "fact_item_sales"."flight_actual_departure_datetime" AS "flight_actual_departure_datetime", "fact_item_sal...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		3	ual_departure_time_id", "fact_item_sales"."flight_id" AS "flight_id", "fact_item_sales"."flight_instance_id" AS "flight_instance_id", "fact_item_sales"."fli...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		4	scheduled_departure_date_id", "fact_item_sales"."flight_scheduled_departure_datetime" AS "flight_scheduled_departure_datetime", "fact_item_sal...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		5	led_departure_time_id", "fact_item_sales"."gate_call_time_id" AS "gate_call_time_id", "fact_item_sales"."gate_number_location_id" AS "gate_number...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		6	"gross_sale_value", "fact_item_sales"."income_to_mag" AS "income_to_mag", "fact_item_sales"."item_quantity" AS "item_quantity", "fact_item_sales"....	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		7	es"."next_destination_airport_id" AS "next_destination_airport_id", "fact_item_sales"."pax" AS "pax", "fact_item_sales"."product_id" AS "product_id",...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		8	, "fact_item_sales"."retail_transaction_date_id" AS "retail_transaction_date_id", "fact_item_sales"."retail_transaction_datetime" AS "retail_transaction...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		9	n_time_id" AS "retail_transaction_time_id", "fact_item_sales"."store_id" AS "store_id", "fact_item_sales"."store_location_id" AS "store_location_id", "fa...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		10	em_sales"."transaction_id" AS "transaction_id", "fact_item_sales"."transaction_item_unique_id" AS "transaction_item_unique_id" FROM "prod"."fact_it...	2017-06-
1009330 fetch 10000 in "SQL_CUR6";		0	close "SQL_CUR6"	2017-06-

Redshift



- Performance so far has been very good.
- A lot to do with the design of Redshift.
 - Optimisations so far have been limited to:
 - Fields:
 - lengths
 - datatypes
 - compression datatypes.
 - Distribution keys.
 - Sort keys.
 - Skew analysis.
 - Vacuum and ANALYZE.
- But we intend to do some more work on below:
 - Work queue management.
 - User load analysis.
 - Attribute pushdown.



Fast

Cost
Efficient

Simple



Elastic



Secure



Compatible

1. Analyze Database **Audit Logs** for Security and Compliance Using Amazon Redshift Spectrum amzn.to/2szR3nf
2. Build a **Healthcare Data Warehouse** Using Amazon EMR, Amazon Redshift, AWS Lambda, and OMOP amzn.to/2rr7LWg
3. Run **Mixed Workloads** with Amazon Redshift Workload Management amzn.to/2srlL1g
4. **Converging Data Silos** to Amazon Redshift Using AWS DMS amzn.to/2klr1bg
5. Powering Amazon Redshift Analytics with Apache **Spark** and **Amazon Machine Learning** amzn.to/2rgR8Z7
6. Using pgpool and Amazon ElastiCache for **Query Caching** with Amazon Redshift amzn.to/2lr66MH
7. Extending Seven Bridges Genomics with Amazon **Redshift and R** amzn.to/2tlylga
8. **Zero Admin** Lambda based Redshift Loader bit.ly/2swvv16

Thank You

Data is magic!

