

# Analyzing Streaming Data in Real-time with Amazon Kinesis

# Agenda

- Why Real-Time Analytics?
- What Is Real-Time Data?
- What Real-Time Services Does AWS Offer?
- Common Use Cases
- Deep Dive

# Why Real-Time Analytics?

# It's All About the Pace

## Batch Processing

---

Hourly server logs

Weekly or monthly bills

Daily web-site clickstream

Daily fraud reports

## Stream Processing

---

Real time metrics

Real time spending alerts/caps

Real time clickstream analysis

Real time detection

# A Day in Life

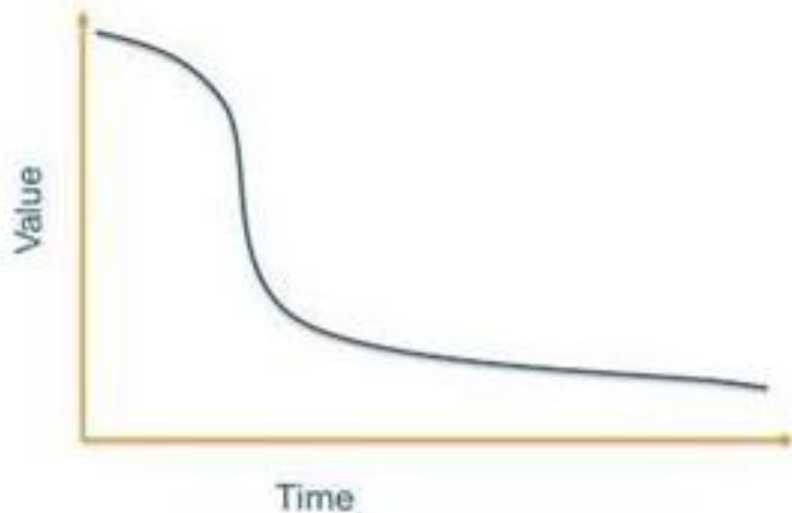


# Data Loses Value Over Time

Ingest data as it is generated

Analyze data in real time to get insights immediately

Deliver data to in seconds instead of hours

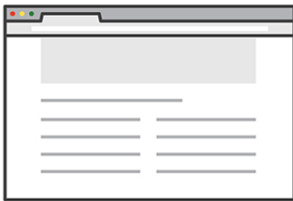


# What Is Real-Time Data?

# What Is Real Time Data?



Mobile Apps



Web Clickstream

```
[Wed Oct 11 14:32:52  
2000] [error] [client  
127.0.0.1] client denied  
by server configuration:  
/export/home/live/ap/htdo  
cs/test
```

Application Logs



Metering Records



IoT Sensors



Smart Buildings



# Simple Pattern for Streaming Data

## Data Producer

Continuously creates data

Continuously writes data to a stream

Can be almost anything



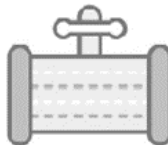
Mobile Client

## Streaming Service

Durably stores data

Provides temporary buffer that preps data

Supports very high-throughput



Stream

## Data Consumer

Continuously processes data

Cleans, prepares, & aggregates

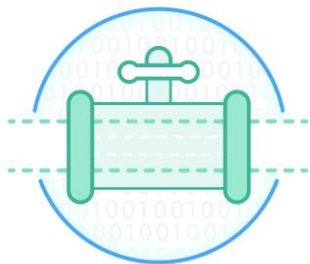
Transforms data to information



Application

# What Real-Time Services Does AWS Offer?

# Amazon Kinesis



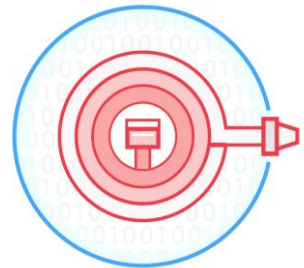
## Amazon Kinesis Data Streams

Build custom  
applications that process  
and analyze streaming  
data



## Amazon Kinesis Data Analytics

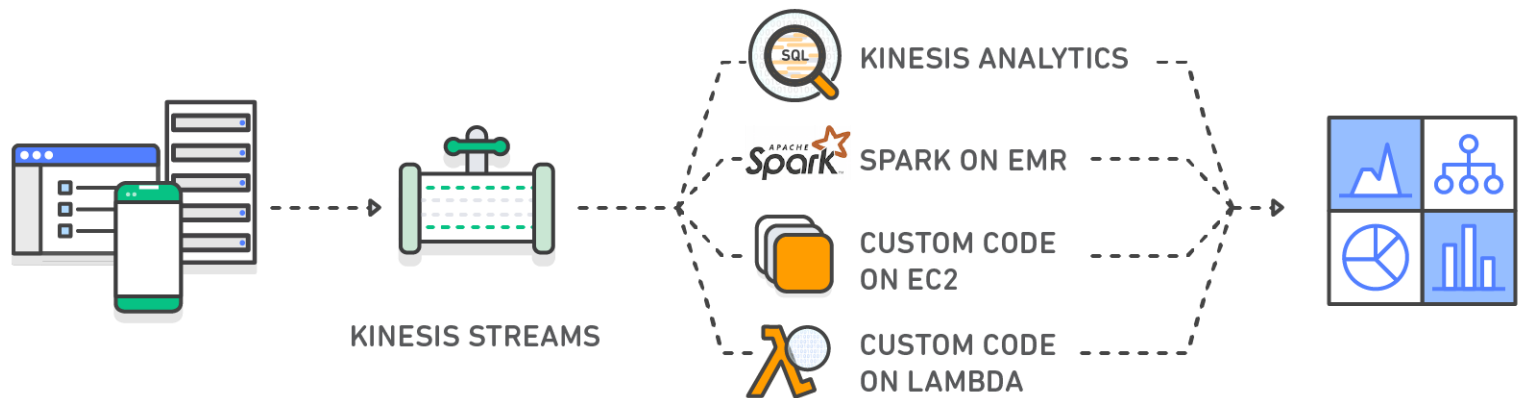
Easily process and  
analyze streaming data  
with standard SQL



## Amazon Kinesis Data Firehose

Easily load streaming  
data into AWS

# Amazon Kinesis Data Streams



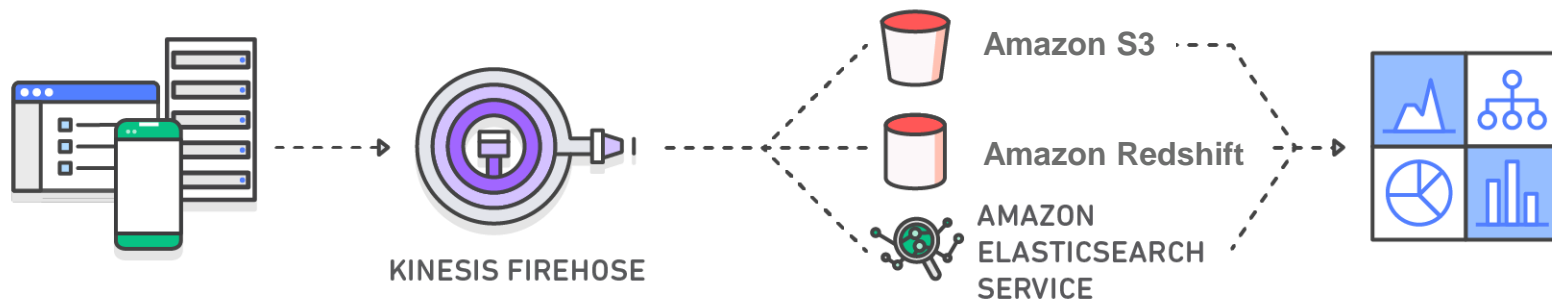
*Capture and send data to Kinesis Streams*

*Build custom, real-time applications using Kinesis Analytics, stream processing frameworks like Apache Spark, or your code running on Amazon EC2 or AWS Lambda*

*Load processed data to any data store, send real-time alerts, feed live dashboards, and more*

- **Easy** administration and low cost
- Build real-time application with framework of choice
- **Secure**, Durable storage

# Amazon Kinesis Data Firehose



*Capture and send data  
to Kinesis Firehose*

*Kinesis Firehose prepares and loads the data continuously to  
the destinations you chose from among S3, Redshift, Amazon  
Elasticsearch Service, and Kinesis Analytics*

*Analyze streaming data using your favorite BI tools*

- **Zero** administration and seamless elasticity
- Direct-to-data store integration
- **Serverless**, continuous data transformation



# Stream Data To Amazon Kinesis

## Automatic ingestion



Amazon  
VPC Flow  
Logs



AWS  
CloudTrail  
Event Logs



Amazon  
CloudWatch  
Logs



AWS IoT  
events



Amazon  
Pinpoint

## Easy setup

As a proxy:



Amazon API  
Gateway



Elastic Load  
Balancing

For change data capture:



Amazon  
DynamoDB



Amazon  
RDS

## Write your own



Amazon  
Kinesis Agent



Amazon  
Kinesis  
Producer  
Library



AWS SDKs

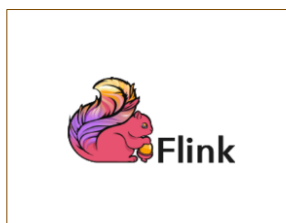
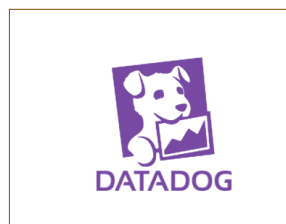
Apache  
**LOG4J**



fluentd

Just a sample... many more ways stream data to Amazon Kinesis

# Integrate With Your Current Solution



# Amazon Kinesis Data Analytics



*Capture streaming data with  
Kinesis Streams or Kinesis Firehose*

*Run standard SQL queries  
against data streams*

*Kinesis Analytics can send processed data to  
analytics tools so you can create alerts and  
respond in real-time*

- **Continuous** anomaly detection
- **Continuous** time series analysis
- **Continuous** filtering
- **Continuous** aggregation
- **Continuous** enrichment



# Amazon Kinesis Data Analytics Applications

100111  
010000  
101001  
010100



Connect to streaming source

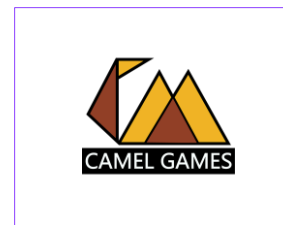
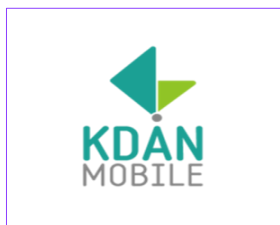
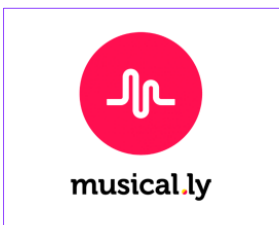


Easily write SQL code to process streaming data



Continuously deliver SQL results

# Amazon Kinesis Customers



# Common Use Cases

# Three Common Scenarios

## **Streaming Ingest- Transform-Load**

Deliver data to analytics tools faster and cheaper

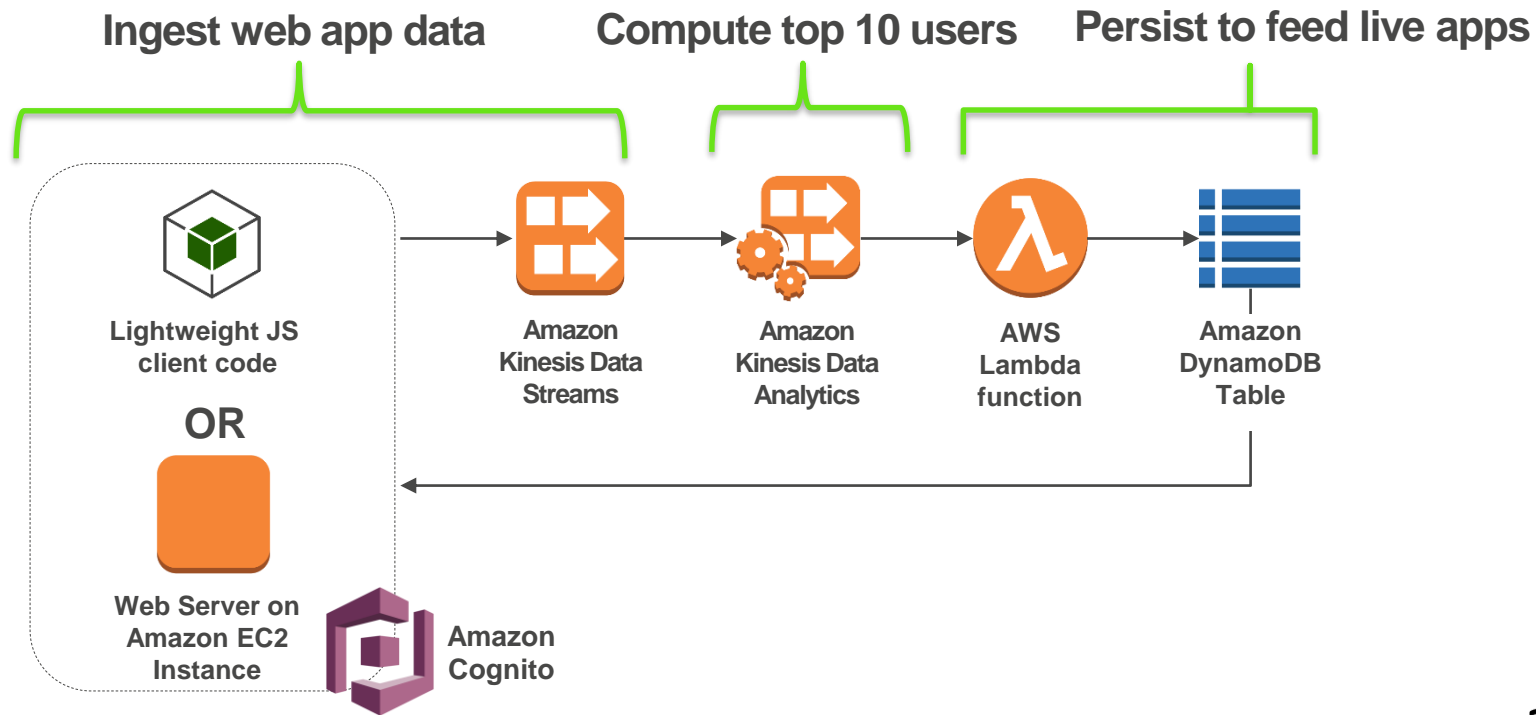
## **Continuous Metric Generation**

Compute analytics as the data is generated

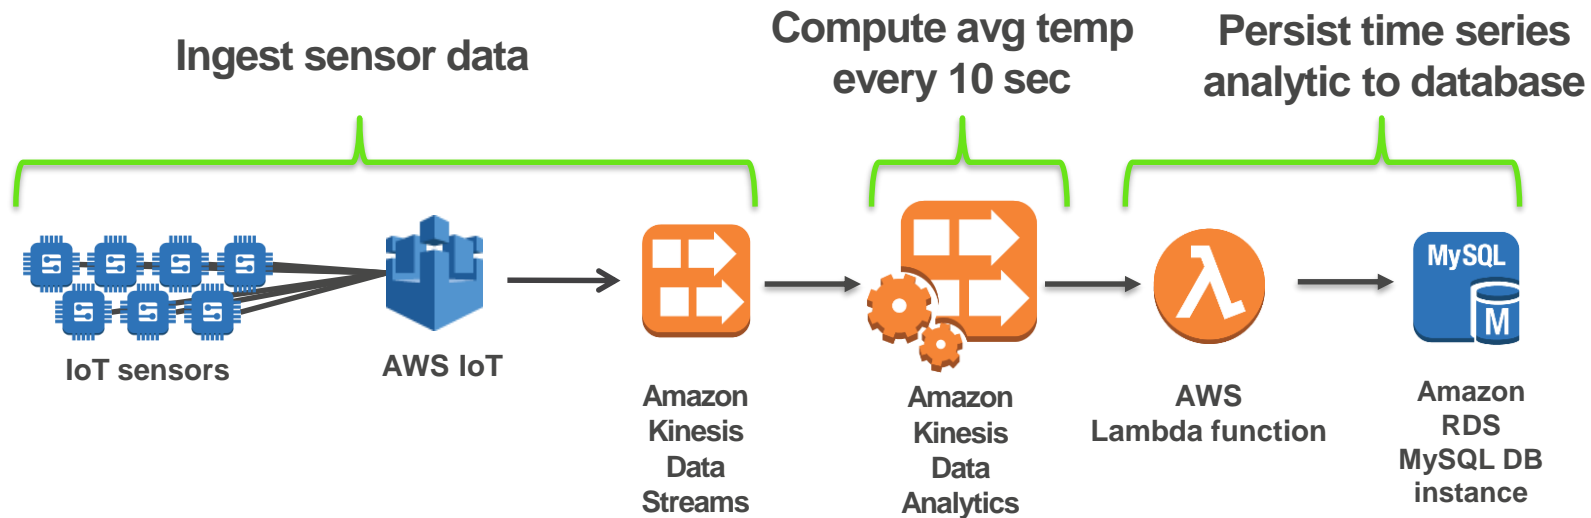
## **Actionable Insights**

React to analytics based off of insights

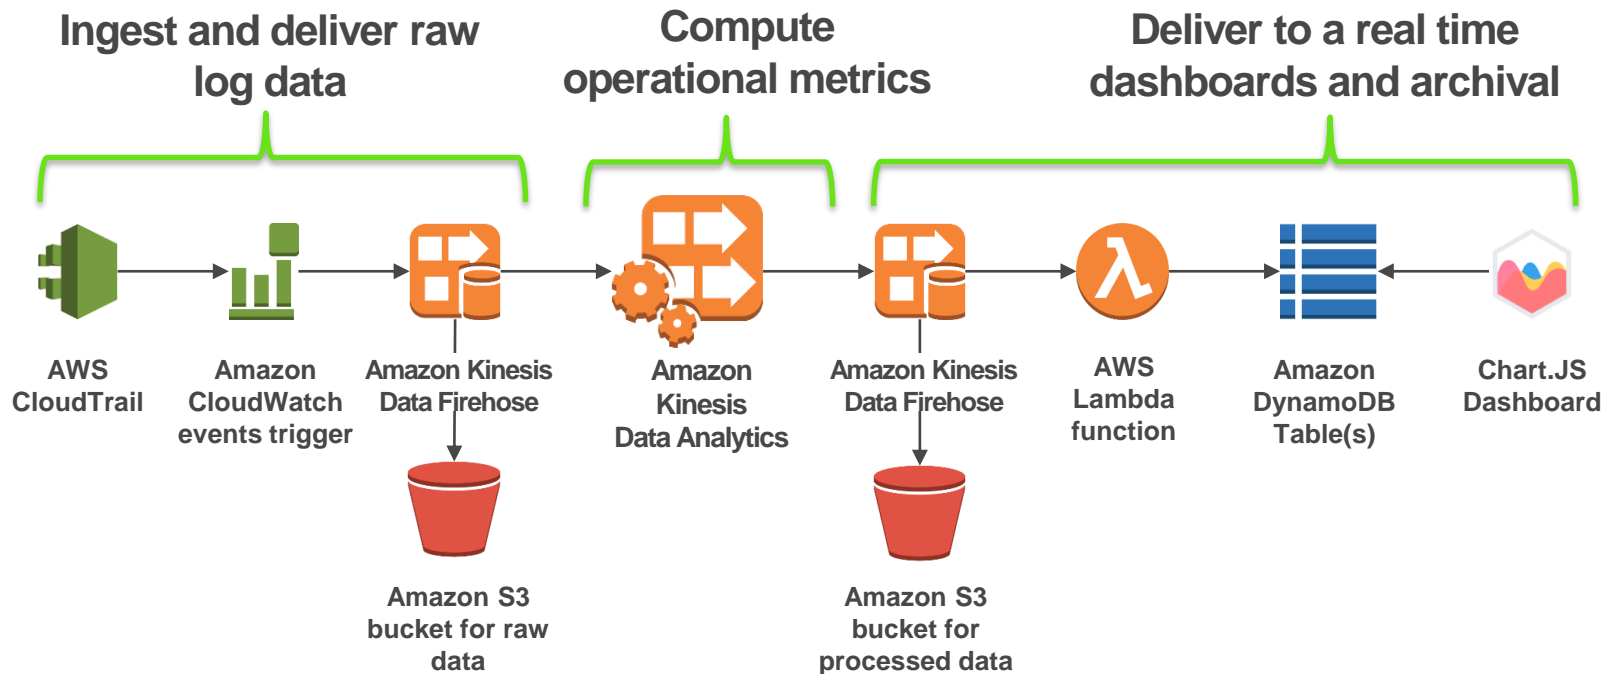
# Web Analytics and Leaderboards



# Monitoring IoT Devices



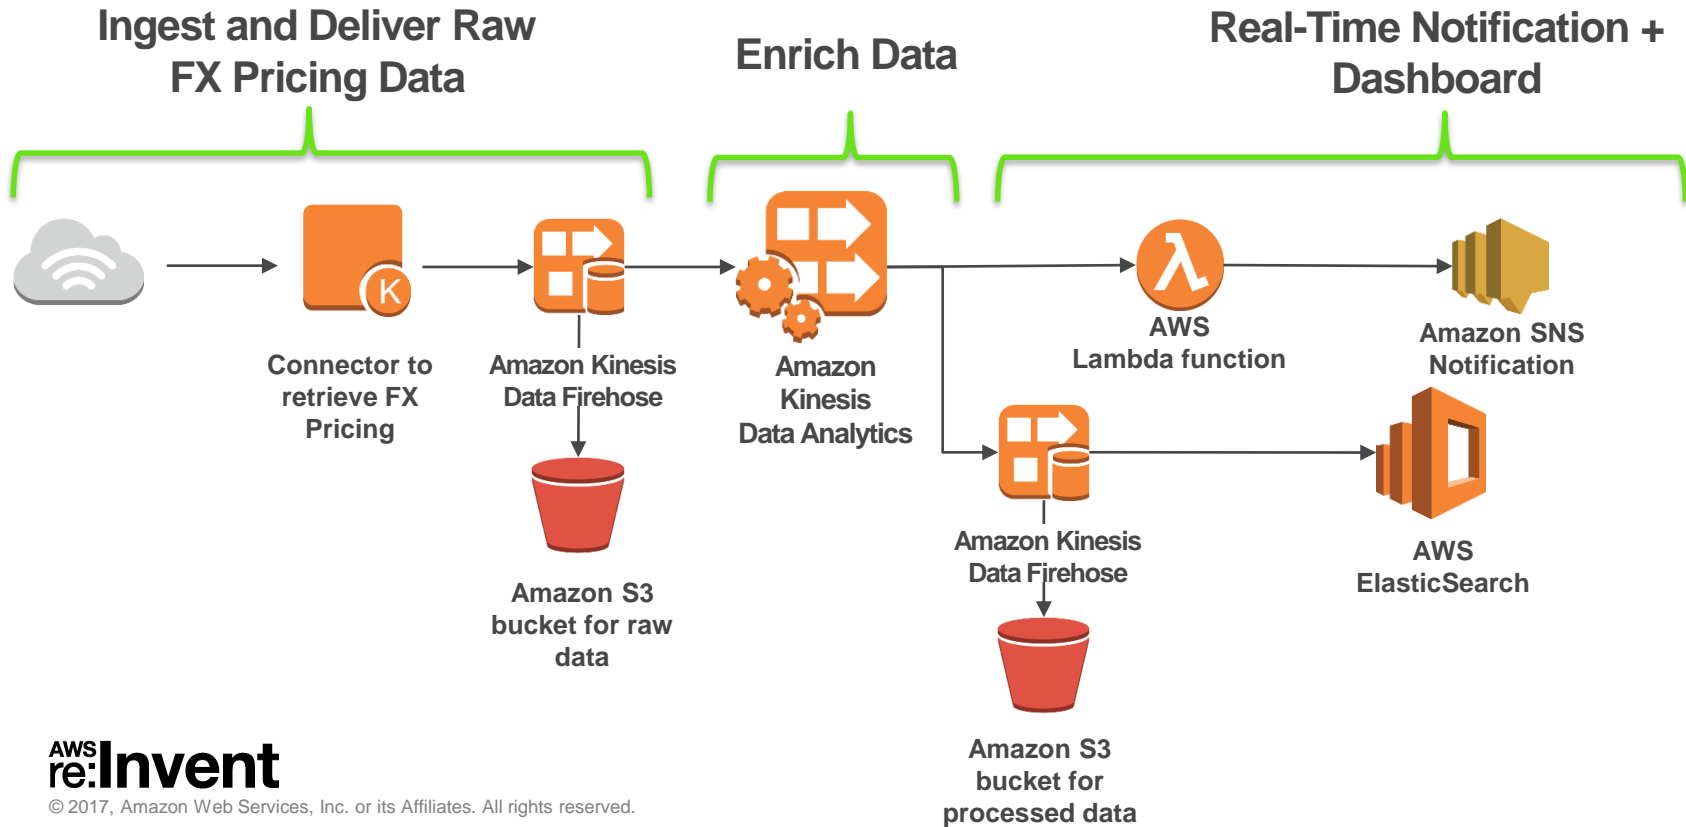
# Analyzing CloudTrail Event Logs



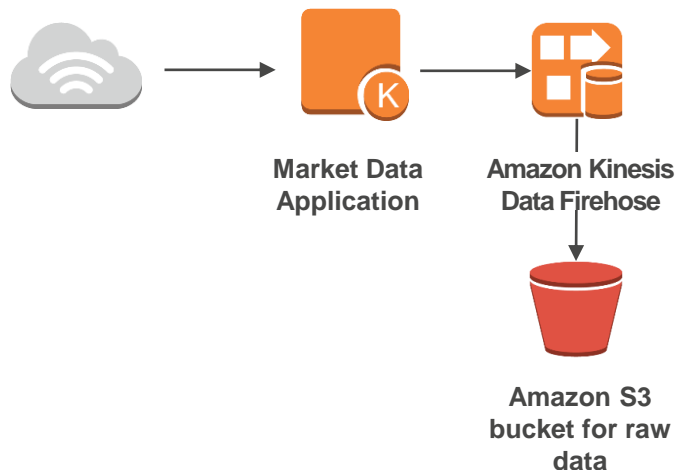
# Analyzing FX Pricing in Near Real Time



# Analyzing FX Pricing in Near Real-Time



# Ingest and Deliver FX Pricing



- An application is connected to a Market Data provider delivering FX pricing, every 30 seconds
- Pricing are sent in real time (to near real time) to Kinesis Data Firehose or Streams
- Each pricing in JSON is in below format

```
{  
  "ask": 212.2553,  
  "timestamp": 1515484223,  
  "symbol": "XAGZAR",  
  "bid": 211.6336,  
  "price": 211.9445  
}
```

# Compute Pricing Metrics in Near Real-Time

Compute metrics using SQL in real time like:

- Max, Min and Average price for the past hour
- Previous Price and Price Change



```
{  
  "min": 212.2553,  
  "timestamp": 1515484223,  
  "symbol": "XAGZAR",  
  "max": 211.6336,  
  "average": 211.9445  
  "change": 1.1954  
  "previous_price ": 210.7491  
  "price ": 211.9445  
}
```

# How Do I Write Streaming SQL? Easy!

Streams (in memory tables)

```
CREATE OR REPLACE STREAM DESTINATION_SQL_STREAM(  
    "@timestamp" BIGINT,  
    symbol VARCHAR(8),  
    min_price DOUBLE,  
    max_price DOUBLE,  
    avg_price DOUBLE,  
    previous_price DOUBLE,  
    new_price DOUBLE  
    change DOUBLE  
);
```

# How Do I Write Streaming SQL? Easy!

## Pumps (continuous query)

```
CREATE OR REPLACE PUMP "STREAM_PUMP" AS
INSERT INTO "DESTINATION_SQL_STREAM"
SELECT STREAM "COL_timestamp" * 1000,
       "ticker_symbol",
       MIN("price") OVER W1 as min_price,
       MAX("price") OVER W1 as max_price,
       AVG("price") OVER W1 as avg_price,
       FIRST_VALUE("price") OVER R1 as previous_price,
       "price" as new_price,
       "price" - FIRST_VALUE("price") OVER R1 as CHANGE
FROM "SOURCE_SQL_STREAM_001"
WINDOW
  W1 AS (PARTITION BY "ticker_symbol" RANGE INTERVAL '1' HOUR PRECEDING),
  R1 AS (PARTITION BY "ticker_symbol" ROWS 1 PRECEDING)
```

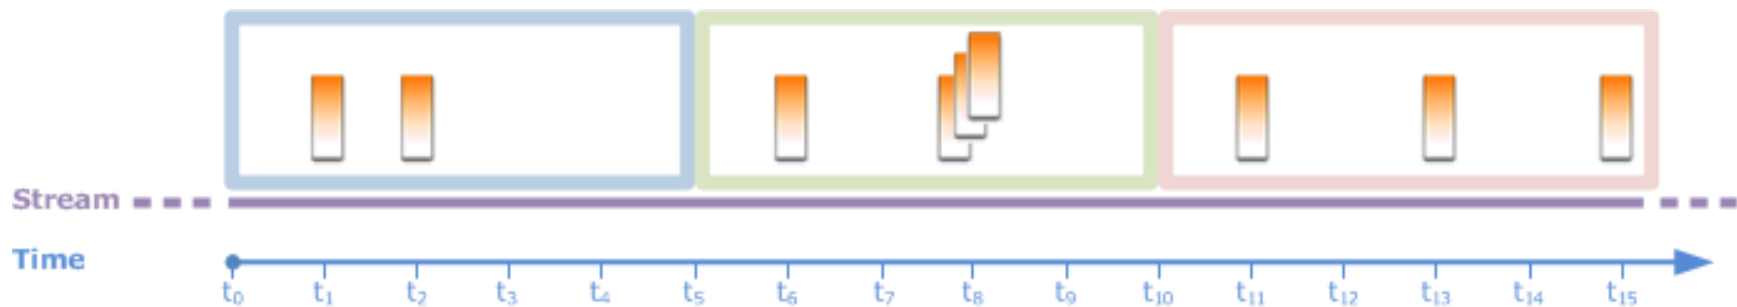
# How Do We Aggregate Streaming Data?

- Aggregations (count, sum, min,...) take granular real time data and turn it into insights
- Data is continuously processed so you need to tell the application when you want results

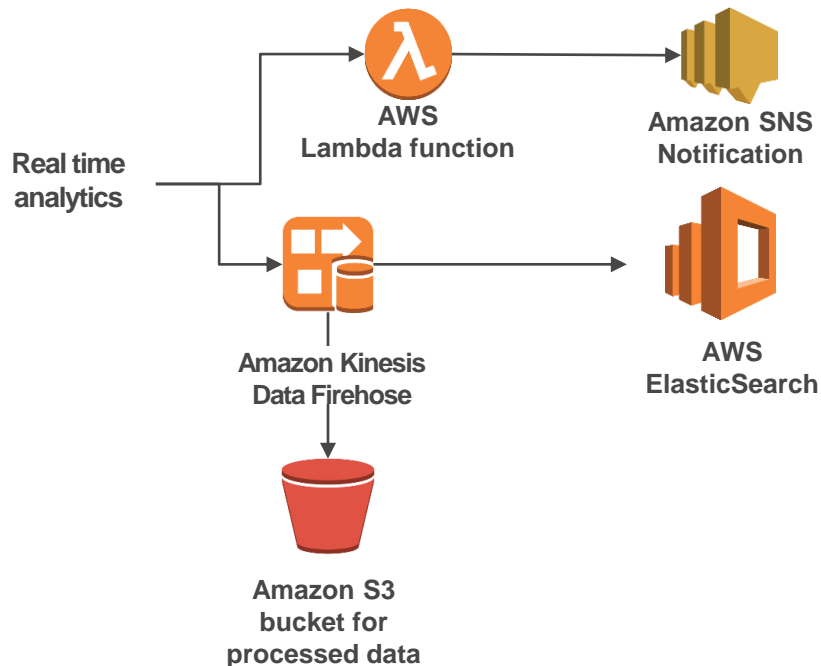
# Windows!

# Window Types

- Sliding, tumbling, and custom windows
- Tumbling windows are fixed size and grouped keys do not overlap



# Real-Time Notification & Dashboard



- Use Kinesis Data Firehose to archive processed data in S3
- Use AWS Lambda to generate notification
- Open source or other tools to visualize the data



# Where To Go Next?

# Getting Started

- Kinesis Home Page: <https://aws.amazon.com/kinesis/>
- Kinesis Blog Post: <https://aws.amazon.com/kinesis/blog-posts/>
- Getting Started Page: <https://aws.amazon.com/kinesis/getting-started/>

# AWS re:Invent

## THANK YOU!

AWS  
re:Invent

